

How Google Serves in the Age of *AI Search*

A senior-level guide to Day 3 in Barcelona: how Google understands queries and ranks pages, what is new in Search Console, and how long its processes take. Every point labelled and checked against Google's documentation.

NOTES & GUIDE BY

Ibrahim Anjro

EVENT

Barcelona · 2 October 2026

EDITION

3 October 2026

Contents

—	Key takeaways	03
	The twelve things to remember from Day 3	
01	Day 3 at a glance	04
	The agenda, a map of the day and the three-day arc	
02	Query understanding	06
	Language, stop words, entities, synonyms, siblings and fan-out queries	
03	Retrieval: from query to candidates	08
	Posting lists; language, country and quality before ranking	
04	What quality means to Google	09
	Ranking systems, the four attributes, testing and raters, Discover	
05	Core updates and spam	12
	Why Search changes, what each update does, and recovery	
06	Search features and rich results	14
	The results page, text results, structured data, product surfaces	
07	Product data for AI shopping	16
	New Merchant Center attributes, product markup, Storebot-Google	
08	Five Search Console launches	18
	Query groups, AI configuration, AI reports, platforms, multimodal	
09	Search Console data and AI visibility	20
	BigQuery and LLMs; showing SEO's value when clicks fall	
10	How AI changes purchase decisions	22
	The messy middle, 'boosted' shoppers, the reassurance step	
11	How long Google's processes take	24
	From URL discovery to recovery after a core update	
12	AI in Search, used responsibly	26
	Hallucinations, deception, AI slop; the three days in one table	
13	Day 3 action checklist	28
	Checks for the content team and the development team	
14	Said on stage, not in Google's docs	29
	Points to treat as informed hints	
15	Help complete this guide	30
	Send what I missed; get the web edition	
—	Sources	31
	Google documentation and press reports used to check this guide	

HOW TO READ THE LABELS

- SLIDE

Shown on screen during the session
- DOCS

Google's published documentation
- ANALYSIS

My interpretation, not Google's words
- STAGE

Said on stage, from my recordings
- PRESS

A third-party report of what Google said
- (not in docs) Said at the event; not in Google's docs

SUMMARY

Key takeaways

Day 3 followed a query from the search box to the results page, then to what site owners can measure. Quality came up at nearly every step: at retrieval, in ranking and in core updates.

01 Google rewrites pretty much every query.

Synonym expansion is traditional Search, not AI. Google knows the synonyms: don't list them on pages.

02 Fan-out queries are ordinary queries.

Google treats them like typed ones. They are not in Search Console; don't chase individual ones.

03 Quality decides retrieval, not just ranking.

The most important signal for ordering candidates, used again in ranking. Weak pages may never compete.

04 Quality means human effort.

Effort, originality, talent or skill, and accuracy. Quality issues are quality issues, not AI versus human.

05 Core updates are not penalties.

No recovery in the penalty sense: improve the site as a whole, and allow several months.

06 AI Overviews and AI Mode are not rich results.

No special markup: a page must be indexed and eligible to be shown with a snippet.

07 Product data feeds AI shopping.

For deeper AI Mode questions: six new Merchant Center attributes. Rich results need only markup; the Shopping tab needs Merchant Center.

08 The generative AI report counts impressions only.

No clicks, no queries; the Performance report still includes AI traffic. Measure visits yourself.

09 Clicks no longer tell the whole story.

Google advises judging the value of visits, such as sales and sign-ups. Report conversions too.

10 AI "boosts" shoppers; it doesn't choose for them.

People hand AI the effort, then double-check on Search and brand sites; familiar brands reassure.

11 Discovery is slow, indexing is fast.

Google's averages: about 20 hours to discover a URL, 30 days to recrawl it, 1.5 hours to index it.

12 AI SEO is SEO.

AI features use exactly the same processes as traditional results, with very few of their own.

01

SECTION 01

Day 3 at a glance



Day 3 followed a query through serving, from query understanding to the results page, then turned to Search Console, measurement and AI. This guide selects from about 700 points; the web edition holds every one.

10:15	● Welcome to serving and ranking day! Google	One slide captured
10:25	● Making sense of users' queries John Mueller and Gary Illyes, Search Relations	Transcript · Slides
10:45	● Lightning session K: Facets of quality Community speakers	Transcript
11:15	● How Google thinks about Quality Google	Transcript · Slides
11:45	● Uncovering Trustworthy Experiences on Discover Google	Slides captured
12:00	● What are quality updates Google	Transcript · Slides
13:10	● Quality prevails	Not covered
13:25	● How Search results are born Gary Illyes, Search Relations	Transcript · Slides
13:40	● Shopping on Search: Beyond the blue links Alex Jansen, Software Engineer, Google Shopping	Transcript
13:55	● Inside Search Console: What's New & How to Use It Ariel Kroszynski, Front-End Engineering Manager, Search Console	Transcript · Slides
14:20	● Lightning session L: Understanding SERPs and your users Nik Vujic and a community speaker	Transcript · One slide
14:45	● Poster session M: Success Strategies That Work Daniel Hernandez Susa, Gaetano Romeo, Victoria Yap, Elif Öztürk, Jo Juliana Turnbull and Erin Sparling	One slide captured
15:30	● Mastering the messy middle Pablo Pérez, Senior Marketing Research & Insights Manager	Transcript · Slides
15:45	● How long does it take to..? Gary Illyes, Search Relations	Transcript · Slides
16:00	● Wrapping all up: AI, Search, and making sense of everything. Gary Illyes, Search Relations	Transcript · One slide
16:20	● Q&A Alex Jansen, John Mueller, Ariel Kroszynski, Duy Nguyen and Gary Illyes, Google panel	One slide captured


17:20

Closing
Google

One slide captured

**AFTER SERVING**

A map of Day 3. The serving steps from Google's opening slide, with the part of this guide that covers each one.



The three-day arc. One stage of Google Search per day. Google said it built the event this way because that is how Search works.

02

SECTION 02 · MAKING SENSE OF USERS' QUERIES · JOHN MUELLER

Query understanding

Google rewrites pretty much every query before it looks anything up: it detects the language, drops stop words, recognises entities and adds synonyms. Fan-out queries from the AI features go through the same steps.

Language, words and entities

STAGE For almost any query, Google first detects the language, which tells it roughly what content the user wants. A brand name alone doesn't reveal it, so Google falls back on location and browser settings and sometimes shows the wrong language version. **DOCS** Search Help adds the Google language setting and the device's languages.

ANALYSIS Don't leave brand searches to chance: connect language versions with hreflang and give every page a visible language switcher.

STAGE Next, Google separates the words, harder in languages such as Thai that put no spaces between them, and drops stop words such as "a" and "of", mirroring indexing. Where small words matter, the whole phrase can be recognised as an entity, and a query naming an entity is read as a request for that thing. Ranking matches words or entities and "probably" mixes both. (not in docs)

Synonyms: pretty much every query is rewritten

STAGE Google's ranking teams "say that the synonym system is one of the most important parts", because users phrase things differently from the way content is indexed. In Mueller's words, "pretty much all queries are rewritten", in traditional Search, not only in the AI features. **DOCS** How Search Works: the synonym system matches "change laptop brightness" with "adjust laptop brightness".



How [fried chicken place in Barcelona] is rewritten. Google's stage example; the real rewrite is internal (not in docs). "Near me" is rewritten the same way, to the user's location rather than the words.

STAGE Synonyms are learned automatically from how people search, not catalogued by linguists: "we don't need to be technically accurate". Words people use interchangeably count as synonyms, which is why a highlighted word in the results can look wrong (not in docs).

RELATION	GOOGLE'S EXAMPLE STAGE	TREATED AS
Synonym	"Photograph" may become "image", "picture" or "photo"; a German look-alike meaning "photographer" was dropped (not in docs)	Interchangeable, per language
Contextual	"GM": General Motors in [GM car], general manager in [GM restaurants], genetically modified in [GM barley] (not in docs)	Depends on the query
Sibling	[Canon camera] should not show Nikon; comparison queries mark such pairs (not in docs)	Related, never swapped
Diacritics	"ü", "ue" or a plain "u"; Google sometimes gets a variant wrong (that part not in docs)	Usually the same
Translation	[Iberico ham] and [jamón ibérico] brought up the same entity (not in docs)	The same thing
Related phrase	[Spanish cured ham] brought up mainly a Wikipedia page at the time (not in docs)	Different: write it out

Don't write for the synonym system

STAGE Google adds synonyms itself, so "you don't have to do that SEO meme of, like, list all the synonyms on your page". It works much as if the user had joined them with OR, so an OR search of a term's synonyms probably returns much the same results (not in docs). "I wouldn't artificially just stuff those variations in there."

DOCS The SEO Starter Guide: anticipate readers' words ("charcuterie" or "cheese board") without worrying about every variation. ANALYSIS Drop synonym blocks; lists of variants read as keyword stuffing.

STAGE Google does not prefer English, and a page needs no English term where Google knows its local equivalent. For diacritics, foreign words and spellings, focus on what users search: Hindi users search in Hindi script and in Latin letters (not in docs).

ANALYSIS · TEST BEFORE YOU WRITE

Search each variant of a key term. Same entity or near-identical results: use your audience's term.

Unrelated results: write that variant on the page. Where users type two scripts or spellings, use the forms Search Console's queries show.

Fan-out queries are queries too

SLIDE "How LLM features with grounding generally work": results go to an LLM, which generates fan-out queries back to the search engine and answers with links. STAGE Fan-out queries are treated like typed ones, synonyms included, and kept distinct from each other for coverage.

STAGE "Fan-out queries are not added in Search Console, because they're basically a part of our infrastructure." (not in docs) They show up in Gemini and some search APIs. DOCS Gemini's grounding responses list the queries run. ANALYSIS Search Console counts AI features under the user's own query.

STAGE "I would not overfocus on which individual fan-out queries happen and how can I rank for those".

ANALYSIS Cover a topic's real subquestions on solid pages instead.

03


SECTION 03 · MAKING SENSE OF USERS' QUERIES, PART 2 · GARY ILLYES

Retrieval: from query to candidates

Retrieval looks the expanded query up in posting lists and already orders the candidates, by language, country and quality, before ranking starts. Google called quality the most important of the three.

STAGE The expanded query goes to the index, which keeps posting lists: for each word, the URLs associated with it (see Day 2). Retrieval matches the query's words and their expansions against those lists (not in docs), using signals attached to each document. The first condition: the query's words, or its concepts in the case of embeddings, are in the document or related to it.

STAGE A query like [best fried chicken ever] matches millions of pages, so Google already orders the candidates during retrieval, before ranking starts (not in docs).



Retrieval as a funnel. The three ordering signals in the order Gary Illyes gave them; the widths are illustrative, not measured.

SIGNAL	WHAT GOOGLE SAID	STAGE
1 · Language	Indexing signals order the candidates, language and country first. A Spanish searcher doesn't necessarily want Italian results; no results in a language is an opportunity.	
2 · Country	A searcher in Switzerland wants Swiss cheese, not German.	DOCS Geotargeting one country can lift rankings there at the expense of other locales.
3 · Quality	The most important: "if the quality of a URL is high, then it's more likely to be retrieved from the index for specific queries." (not in docs) The same signal is used again in ranking.	

ANALYSIS · QUALITY ACTS TWICE

A low-quality page is less likely to be retrieved at all, so it may never reach ranking: quality is a precondition, not a final tweak. The docs list quality among ranking signals but never say it decides retrieval. Language and country act just as early, so make both unambiguous: one main language per page, hreflang between versions, local prices, addresses and shipping.

04

SECTION 04 · HOW GOOGLE THINKS ABOUT QUALITY · LIGHTNING SESSION K · DISCOVER

What quality means to Google

Google defined quality as the effort, originality, talent or skill and accuracy a person put into content, and called it one of the most important of hundreds of ranking signals. A community talk argued that SEO is for humans, and the Discover talk put trust first.

Not one system, but hundreds of signals



"There's not one single ranking system." The systems named on Google's slide: examples, not the full set.


STAGE Quality sits in the core ranking systems as a signal of its own, a number calculated from many parts (not in docs): one of the most important, but one among hundreds. "Big takeaway: always remember, there's hundreds of ranking signals." For news, freshness, diversity and originality probably weigh more. **ANALYSIS** Improve content, links, page experience and freshness together, not one metric.

STAGE PageRank, the speaker's favourite, "even though we don't really use them so much anymore" (not in docs). **DOCS** It has evolved and is still a core system. **ANALYSIS** It weighs less, but links still matter. And per the guide MUM is not used for general ranking: the slide lists systems Google runs.

Quality, defined: four attributes


"Quality is basically the extent to which a human being put effort, originality, talent, skill, and accuracy into creating content."

Google's quality talk, citing page 21 of the rater guidelines



DOCS Google's helpful-content guidance lists the same four attributes, and says generative AI used to mass-produce text without oversight represents little to no effort. **ANALYSIS** Google updated that page on 1 October 2026, during the event, to match the talk: cite it, not the talk. The rater guidelines ask whether a page is the original source: aim for that, and add substantial value where you build on others.

How a change reaches Search



From idea to launch. The steps described in the quality talk; the 0.1% start and the raters' "which and why" are from Google's How Search Works page.

719,326	124,942	16,871	4,781
search quality tests	side-by-side experiments	live traffic experiments	launches to Search
2023, Google's slide and How Search Works	2023, same sources	2023, same sources	2023, same sources

STAGE Thousands of experiments are probably running at any moment. **ANALYSIS** A later slide rounded the figures to 800,000+ tests and 4,700+ launches; cite the exact ones, with the year. Against nine named ranking updates on the status dashboard in 2023, almost every launch goes unannounced: monitor continuously, not only around announced updates.

Search quality raters

SLIDE Independent reviewers who help Google judge a change before launch, working from the Search Quality Rater Guidelines (goo.gle/raters-guidelines). "They can not affect individual site's rankings." **STAGE** They are external, and may also work for other search engines (that last point not in docs). They rate results for a change and explain their rating.

STAGE Ratings measure how well search engines deliver helpful content; nothing in the guidelines is itself a ranking factor. **ANALYSIS** Read them as what Google's systems aim to reward, not as a checklist.

Commodity and non-commodity content

	GOOGLE'S QUALITY TALK STAGE	GOOGLE'S AI FEATURES GUIDE DOCS
Commodity	"Top 10 things to consider when buying running shoes": common knowledge anyone could write, little insight.	"7 Tips for First-Time Homebuyers".
Non-commodity	Why certain running shoes collapse after 400 miles, from wear-pattern analysis: a unique, experienced take.	A first-hand account of waiving a home inspection.
Strong example	A review site that cut shoes apart, studied the foam under a microscope, built a test rig and its own metrics (not named).	Reviews earn trust by showing how many products were tested, the results and the method.

STAGE Illustrations, not a recipe. On AI content, the speaker gave a personal opinion and one firm line: "it's very important to think about these quality issues as quality issues and not AI versus human content."
ANALYSIS Ask whether an article adds more than an LLM would answer directly. For reviews: test the products, publish your own photos and measurements, explain the method.

Lightning session K: SEO is for humans

STAGE A community speaker: "SEO is for humans". Around 2009 writers bent titles to hit their CMS's keyword-density score, and it worked at the time (not in docs). **DOCS** Today that is keyword stuffing, and Google has no preferred word count. **ANALYSIS** Plug-in scores are not editorial targets.


STAGE Entering Poland late with a small team, a large Spanish software-download site went deep instead of broad, and its Polish edition reached the top three Polish technology sites and double the organic traffic of its English edition (period not stated).

A COMMUNITY CASE STUDY: NICHE BEATS BROAD

STAGE A complete, well-written article got 1,600,000 impressions; one for a very specific audience (friends looking for an alternative to a failing game server) got 46,000 in the same period, yet 28 times more clicks, and its visitors converted 10 times more. **ANALYSIS** No source or absolute numbers were given, so treat it as an illustration, not a benchmark.

STAGE "walk out of your shop. Stand on the other side of the street, look at the shop the way your users are looking at it". If you can't say what problems your pages solve, you are optimising for the search engine.

Discover: trust first



The Discover gate. Redrawn from Google's Discover slide. Google's docs add that indexed content meeting the content policies is eligible automatically, with no special tags or markup.

DOCS Discover uses many of the same signals and systems as Search. **SLIDE** E-E-A-T: first-hand experience, the knowledge or skill, and being a go-to source; and "untrustworthy pages will have low E-E-A-T no matter how Experienced, Expert or Authoritative they may seem".

SLIDE Images at least 1,200 px wide, over 300,000 pixels, 16×9, enabled by `max-image-preview:large`; no site logos, text-heavy or misleading images. **DOCS** Name the image in schema.org markup or `og:image`; news sources need clear dates, bylines, author and publisher details and contact information.

ANALYSIS On article templates: `max-image-preview:large`, a relevant 1,200 px image in `og:image`, bylines and dates.

05

SECTION 05 · WHAT ARE QUALITY UPDATES

Core updates and spam

Google changes Search thousands of times a year. Core updates reassess quality across the web and punish no one; spam updates target spam only, and only they come with a recovery in the usual sense.

Why Search changes

SLIDE Three reasons: new content formats, which may get Search features once enough people look for them (not in docs); a broader web, where the best content is harder to find (core updates); and spammers finding loopholes (targeted algorithms).

STAGE The speaker's own label for feature launches is "feature updates": "It's not an official name." Every update has one goal: better results. **DOCS** Google's first index held 26 million pages in 1998. **ANALYSIS** Quote that, not the rounded 25 million said on stage. **STAGE** For [durian], each result of 2000 faces about 20,000 in 2025 (not in docs).

40 billion

spammy pages Google discovers every day; on stage, "tens of billions"

Google's published figure (webspam report for 2020); the slide put it under "In 2023"

>99%

of visits from Search spam-free: Google's aim, which it says it meets

Measured as visits or searches, not results

2 + 4

core updates (March, May) and spam updates (March, June, August, September) in 2026

Search Status Dashboard, 3 October 2026

Core update or spam update?

	CORE UPDATE	SPAM UPDATE
What it is	STAGE A broad change to Google's systems, several times a year, posted on the Search Status Dashboard.	DOCS A notable improvement to the always-running spam systems, such as SpamBrain. The September 2026 one began on 24 September.
Targets	STAGE No sites: it changes how Google assesses properties of content. Pages that drop usually break no spam policy.	STAGE Spam only: new loopholes, and sites churning out millions of low-quality pages for money.
"Recovery"?	STAGE "Core updates are not punitive." Technically none, "because you were not penalized".	SLIDE Yes: remove the kind of spam the update's announcement describes.
How long	DOCS Days to months, sometimes the next core update; smaller, unannounced core updates can also help.	DOCS Months. After a link spam update, the benefit of the spammy links is gone for good.

STAGE When pages drop, others rise: a core update may reward under-rewarded pages. It sometimes comes with new rater guidelines (not in docs). **DOCS** Each is tested first, with rater feedback not used directly in ranking.

After a core update drop

STAGE Look at the whole site: "you should check if your site overall, not just individual pages, is delivering content that's helpful, reliable, and people-first". Ask someone else to review it, and ask users, for example with CSAT surveys (not in docs).

DOCS Wait a full week after the rollout ends, then compare with a week before it. A drop from 2 to 4 needs no drastic action; 4 to 29 needs a deeper look. **STAGE** Improving the whole site can win rankings back, but a competitor may simply be better now. **SLIDE** Learn from the sites doing better.

ANALYSIS · A REVIEW PLAN, NOT A "RECOVERY"

ANALYSIS Plan a site-wide quality review instead of page-by-page fixes: an outside reviewer, a comparison with the pages that gained, and a test of whether users can finish their task.

Spam: the five types on the slide, and what replaced them

STAGE Spam is content designed to deceive users or manipulate search results. **SLIDE** "Spam, spam and spam!" listed five types (more at goo.gle/spam-policy).

TYPE	GOOGLE'S SLIDE SLIDE	CHECK FOR DEVELOPERS ANALYSIS
Cloaking	Different content for users and search engines, to manipulate rankings.	Serve Googlebot what users get.
Doorways	Pages made to rank for similar queries, leading to less useful pages.	No near-identical location or keyword pages.
Scraped content	Content taken from other sites without adding value.	Add value to any reused content.
Link spam	Any links meant to manipulate rankings.	Today, audit mass-produced pages first.
Hacked content	Content placed through security holes; prevention is key.	Patch and monitor the CMS for injected pages.

STAGE "Some of these, I would say, are not that common anymore, or we don't care all that much about them." Scaled content abuse would replace link spam (both points not in docs). AI lets Google test more candidates, launch spam updates faster and catch new spam; four had shipped in 2026 by 2 October. **ANALYSIS** Audit mass-produced pages, AI-made ones included, before old link patterns.

NEW POLICY

Back button hijacking

STAGE Added after many user complaints. **DOCS** Announced 13 April 2026, enforced from 15 June 2026; ad or library code can cause it. **ANALYSIS** Test that back returns users to the results; audit History API and vendor scripts.

CLARIFIED

Manipulating AI responses

STAGE Spam aimed at AI answers breaks the policies too. **ANALYSIS** Hidden text for AI or pages for fan-out queries: same rules. **STAGE** Spam reports steer Google's work. **DOCS** Report text goes to the site owner.

06

SECTION 06 · HOW SEARCH RESULTS ARE BORN · GARY ILLYES

Search features and rich results

Search features are built as the last serving step, after ranking, and most need nothing from the site owner. Rich results are the exception, usually built from structured data; AI Overviews and AI Mode are not rich results.


STAGE Gary Illyes called the results page an auction where result types bid for space (a metaphor, not ads). Text results are the most common type (not in docs).

Query understanding picks the result types

STAGE Query understanding predicts intent, and with it the result types. [orange pictures] names the type; for an ambiguous [orange], Google looks at how users reacted to result types for similar queries. **DOCS** Aggregated, anonymized interaction data becomes signals for its machine-learned systems.

STAGE From 2005–2006 people added "pictures" to queries, which led to universal search (not in docs). **DOCS** Launched 16 May 2007. **ANALYSIS** The launch post cites other reasons. **STAGE** An image query is broken into embeddings searched for in the index, and its intent can still be detected (not in docs).

Anatomy of a text result



A text result, drawn abstractly as Google's Visual Elements gallery does, so it doesn't date like a screenshot. Illustrative page and site.

STAGE Title links and snippets come from Google's understanding of the page, even when the owner adds nothing. Image results come from the image index (not in docs); video thumbnails and attribution from the page or its markup.

Search features come last; rich results need data

SLIDE Search Features is the last step on Google's serving diagram. **STAGE** Results come back from the index and are ranked; then the features are built, mostly from what Google extracted at indexing, with nothing needed from the owner. Rich results are the exception: "The magic behind rich results is usually structured data. Emphasis on 'usually'." Usually JSON-LD.

AI FEATURES ARE NOT RICH RESULTS

STAGE "AI Mode and AI Overviews are not rich results." They work from the normal text results; only future parts such as images or recipes would need structured data. **DOCS** No special schema.org markup: a page must be indexed and eligible for a snippet.

ANALYSIS Day 2 said processed markup also feeds these features: markup is no condition for appearing, but its data can still reach them.

Articles, site names and reviews

SLIDE Article markup: "Not required for 'Top Stories,' but highly recommended for all articles." It can also flag paywalled content. The slide showed the `NewsArticle` example from Google's docs (redrawn here). **DOCS** It tells Google explicitly that the page is a news article, its author and its title.

STAGE The site name is an attribution feature, not a rich result. **DOCS** Google generates it; state a preferred name with `WebSite` markup on the home page.

```
<!-- in <head>, example values -->
<script type="application/ld+json">
{ "@context": "https://schema.org",
  "@type": "NewsArticle",
  "headline": "Running shoes, tested",
  "image": [ "https://example.com/1x1.jpg",
             "https://example.com/4x3.jpg",
             "https://example.com/16x9.jpg" ],
  "datePublished": "2026-10-02T08:00+02:00",
  "dateModified": "2026-10-03T09:20+02:00",
  "author": [{ "@type": "Person",
                "name": "Jane Doe",
                "url": "https://example.com/jane" }] }
</script>
```

STAGE Review snippets show the average rating and often the review count, for products, businesses, films or books: "Those little yellow stars are a powerful signal of quality and trust for our users" (not in docs). **DOCS** Self-serving reviews of your own business are not eligible.

Which product surface needs what

SURFACE	PRODUCT MARKUP	MERCHANT CENTER	NOTES
Product rich results	✓ Enough	Optional	STAGE "Essential" for e-commerce. DOCS Both together maximise eligibility.
Google Images product annotations	✓ Enough	Not needed	STAGE Also built from product markup.
Shopping tab	✗ Not enough	✓ Required	STAGE Submit products in Merchant Center. DOCS Crawled products can appear on Google without either; the page doesn't say if that covers this tab. ANALYSIS Assume a feed.
Feed checks	Helps	Doesn't need markup	STAGE Markup can help validate. DOCS Automatic item updates use it to fix price and availability. ANALYSIS Keep both identical.

07

SECTION 07 · SHOPPING ON SEARCH · ALEX JANSEN, GOOGLE SHOPPING

Product data for AI shopping

AI Mode answers deeper product questions than classic shopping results, so Google added a few flexible Merchant Center attributes rather than hundreds of narrow ones. Product data quality matters more than ever, for AI and for agents that buy on a user's behalf.

STAGE Google Shopping gets billions of product offers from Merchant Center feeds and from schema.org product markup on websites (that part not in docs); both power product snippets in Search. Classic results have little room, so a few attributes were enough: title, image, price, perhaps ratings, shipping and returns. Behind them, brand and barcode (UPC) drive query understanding, retrieval and ranking, and still matter as much.

STAGE AI Mode must return far more detail, as shoppers keep asking deeper follow-up questions. Data quality matters most for agents that act for a consumer: "If the agent has bad data, it cannot do the right things." It might buy the wrong product (the agent point not in docs).

Six new attributes, not hundreds

STAGE After merchants asked whether AI experiences represented them correctly, Google added feed attributes in 2026, announced at Google Marketing Live. It rejected hundreds of thousands of narrow ones, such as heel height or lens details: hard work for merchants and no longer needed for LLMs (not in docs). **DOCS** Google announced "dozens of new data attributes" on 11 January 2026 for a small group of retailers, and opened conversational attributes to retailers globally on 20 May 2026.

FEED ATTRIBUTE	WHAT IT CARRIES STAGE	DOCUMENTED LIMITS DOCS
<code>question_and_answer</code>	Q&As from the maker or consumers, such as a kettle's temperature settings	Up to 30 pairs; 1,000 characters each, 10,000 in total; no prices, shipping, dates or company name
<code>document_link</code>	A PDF manual, spec sheet or package contents	Up to five PDF URLs; the merchant must own the licensing rights
<code>related_product</code>	Accessories, parts and replacements, such as what to buy with a printer	With the two above and <code>popularity_rank</code> : mainly for AI Mode
<code>item_group_title</code> , <code>variant_option</code>	Variants described far more flexibly than colour and size	ANALYSIS The only new attribute the help maps to schema.org: <code>item_group_title</code> to <code>ProductGroup.name</code>
<code>popularity_rank</code>	How well the product sells on a marketplace	0–100, set by the merchant from recent sales against its own inventory. ANALYSIS Relative and self-reported, not a sales figure
<code>product_highlight</code> (existing)	A short bulleted list, now said to help AI systems	2 to 100 highlights of up to 150 characters; no promotional text or keywords
<code>product_detail</code> (existing)	Specifications as key-value pairs, instead of one attribute per spec	Up to 100, each a section, attribute name and value

ANALYSIS All six new attributes (the first five rows) are in the Merchant Center specification on 3 October 2026; January's "dozens" was the wider count. **STAGE** "rich, lightly structured information is critical to AI systems": some structure, no deep nesting, and rich data no standard prescribes, because the merchant knows best (not in docs).

Feeds first; schema.org follows

STAGE Most new attributes were said to exist in schema.org already, so Google can read them from markup too, and a recent schema.org release added `specification` and `valueGroup` for product details (not in docs). Google works to make what feeds and the Universal Commerce Protocol (UCP) express expressible in schema.org. **DOCS** UCP, launched on 11 January 2026, is an open standard for agentic commerce, compatible with A2A, AP2 and MCP.

THE FEED IS THE DOCUMENTED ROUTE

ANALYSIS On 3 October 2026 neither the Merchant Center help nor Search Central documented schema.org markup for the other new attributes, and schema.org's release 30.1 (16 September 2026) is not Google documentation. Send them in the feed.

ANALYSIS **Developers:** keep price, availability, brand and GTIN accurate first; then add Q&As, highlights, key-value product details and related products in the feed, and mirror them in markup only as Google documents support. **STAGE** Google called this early days.

ANALYSIS **Content teams:** fill the Q&A attribute from real customer questions (support tickets, reviews, sales chats), answered factually per product.

Product markup: shared policies by `@id`

STAGE Product markup has product data that is the same for every merchant (name, description, brand, images) and offer data for one merchant (URL, condition, availability, price); an offer can point to shipping and return data defined elsewhere. **DOCS** Define shipping and return policies once under Organization and reference them from each Offer, even on another page, with only the `@id` keyword: through `hasShippingService` (inside `shippingDetails`) for shipping, `hasMerchantReturnPolicy` for returns.

```
// Once, under Organization
"hasMerchantReturnPolicy": {
  "@type": "MerchantReturnPolicy",
  "@id": "https://example.com/#return",
  ... },
// In every Offer: only the @id
"hasMerchantReturnPolicy": {
  "@id": "https://example.com/#return"
}
```

SHOPPING'S OWN CRAWLER

Storebot-Google

STAGE Separate from Googlebot: it needs fresh prices and availability, crawls far more often, checks feed prices on the site and deep-crawls checkout pages. **DOCS** It can fill in checkout forms; robots.txt rules for it affect all Shopping surfaces. **ANALYSIS** Don't block it on product, cart or checkout paths, in robots.txt or bot protection.

50%

of the time, lululemon's own conversational attributes were used in relevant AI Mode recommendations


Google's testing with one brand; Google blog, 16 September 2026

08

SECTION 08 · INSIDE SEARCH CONSOLE · ARIEL KROSZYNSKI

Five Search Console launches

Google presented five Search Console launches of the past year: query groups, AI-powered configuration of the Performance report, generative AI control and reporting, platform properties for social accounts, and multimodal search reporting.



A year of launches. Dates of Google's announcements. The AI-powered configuration was presented without a date and is still shown to only a small share of users.

LAUNCH	WHAT IT SHOWS STAGE	WHERE	DOCUMENTED LIMITS DOCS
Query groups	Queries with the same intent in one group, titled by its most prominent query, with the group's total clicks	Insights report; a click opens the Performance report	Computed with AI; groups may change; no effect on ranking; only for properties with many queries
AI-powered configuration	A chat box that turns a plain-language request into Performance report settings	Performance report for Search results	Experimental; a small share of users; 20 requests per user per day
Generative AI reports	Impressions in Google's AI surfaces, by page, country and device	Left navigation, under Search results	Impressions only: no clicks, no query dimension
Platform properties	Google traffic to an Instagram, TikTok, X or YouTube account	A property per account, verified through the platform	Google Search, Discover and News only; ownership rechecked; data after a few days
Multimodal search reporting	Searches with Lens, Circle to Search, uploaded images and Chrome's image search	Performance report, multimodal search type	Also in the generative AI report; exportable; global from 24 September 2026

1 · Query groups

STAGE Query groups put together queries with the same intent in different words: more or fewer words, or spelling errors. **SLIDE** "How to spell Britney Spears?" **STAGE** Google's logs show hundreds of ways to write it. Grouping goes beyond matching text; granularity is the hard part, and some queries aren't understood or land in illogical groups (that part not in docs).

STAGE Groups follow meaning: "Gemini" meant the star sign in 2023 and may be grouped differently now. A click on a group opens the Performance report with a regex filter of all its queries. **ANALYSIS** The archive page on the slide lists 593 spellings (about 2011). The launch was placed in December on stage; Google's post is dated 27 October 2025. Read groups as intent themes: write for the intent, not each spelling.

2 · AI-powered configuration

STAGE The Performance report is complex, so users now describe the analysis in a chat box, in any language (not in docs), for example: "Show me queries on phone searches that contain the word sports in the last 6 months." It lists its changes and applies them after approval; follow-ups build on them: "So it's not one-shot, but you can also interact with it." In Google's test, "wine-related queries that lead to my site" produced a regex of wine names, such as Merlot, joined by pipes.

CHECK THE AI-BUILT FILTER

DOCS It keeps existing report filters when it adds new ones. **ANALYSIS** Read the proposed filters and regexes before trusting the numbers: the feature is experimental, AI can misread requests, and it covers only the Search results report, without sorting or export.

3 · Generative AI control and reporting

STAGE Every property owner can tell Google whether the site's content may be used in AI search results. A report under Search results shows impressions in Google's AI surfaces by page, country and device, and should get richer; the regular Performance report still includes AI traffic. **ANALYSIS** No clicks, no queries: measure the visits that follow AI answers in your own analytics.

4 · Platform properties

STAGE Launched three or four months before the event, they bring Instagram, TikTok, X and YouTube accounts into Search Console, with more platforms to come (that part not in docs). Instead of a token on the homepage or in DNS, a pop-up from the platform asks the user to let Google identify the account, which proves ownership. A landing page shows clicks, search types, top content, trends, query groups and countries, and Google published a guide to using them.

```
# Performance > Compare > Page
URLs containing /watch
# YouTube long-form
URLs containing /shorts/
# YouTube Shorts
# Instagram: /p/ vs reels
```

STAGE In the example channel, long-form videos got more Google traffic than Shorts: one channel, not a rule (not in docs). **DOCS** One property per account; access pauses when the platform login expires; the data covers Google Search, Discover and News only; accounts of a claimed Search profile are added automatically.

ANALYSIS Media teams: add every account, compare formats by URL pattern.

5 · Multimodal search reporting

SLIDE The newest launch, under a "Hot from the oven" badge. **STAGE** Until about a week and a half before the talk, the Performance report showed essentially text searches. Now Google Lens, Circle to Search, uploaded images and Chrome's right-click image search are counted under the multimodal search type. **DOCS** Announced 24 September 2026; also in the generative AI report, and exportable.

ANALYSIS This search type has no queries dimension: judge image-led traffic by page, country and device, and make product images distinctive enough for Lens to match. **STAGE** The team reads the thumbs-up and thumbs-down feedback on every new feature.

09


SECTION 09 · LIGHTNING SESSION L: UNDERSTANDING SERPS AND YOUR USERS

Search Console data and AI visibility

Two community talks. One feeds Search Console's bulk export to an LLM to support decisions for clients; the other argues that rankings and clicks no longer show SEO's value once AI answers take part of the demand.

Search Console data, fed to an LLM

STAGE Nik Vujic, an agency founder, feeds Search Console and GA4 data to an LLM agent to support client decisions. Search Console is his favourite SEO tool; third-party tools only estimate. The workflow starts with the bulk export to BigQuery.



The community workflow. Three steps from Nik Vujic's slides, which add: "No agent needed to start: pull the data and import it into the LLM of your choice, and talk with it."

SLIDE "Apply the filters you need, even natively in GSC. For example branded vs non-branded, with clicks and impressions." **DOCS** The branded filter is an AI-assisted classification, not a regular expression, and may misidentify some queries. It works only for top-level properties with enough queries and impressions, and has been open to all eligible sites since 11 March 2026.

ROUTE	LIMIT	WHAT TO KNOW
Interface export	1,000 rows	SLIDE "The UI export is limited to 1,000 rows." DOCS Confirmed.
Search Analytics API, Looker Studio	50,000 rows per day per site per search type	STAGE Looker Studio was the second way past the limits. ANALYSIS It lifts the 1,000-row limit, not the daily cap.
Bulk export to BigQuery	No daily row limit	ANALYSIS "Anonymized queries excluded" means the text: rows stay with an empty query, often the top "query". Have the LLM skip them in rankings but keep them in totals.
Any filtered view	Anonymized queries left out	DOCS Filtered rows don't add up to the chart totals. ANALYSIS Branded plus non-branded won't match the unfiltered total; tell the LLM the gap is expected.

ANALYSIS Spot-check queries in each branded group before an LLM sees the split; on URL-path or subdomain properties, use a regex query filter.

Uses, and where the data misleads

- STAGE** Compare articles; to report a refresh on a large site, let the LLM count the queries the content improved on.
- STAGE** After publishing, queries show up first, clicks and impressions much later (not in docs).
ANALYSIS A query only shows through an impression: what rises first is the number of queries with a few impressions each.

- **STAGE** Refresh, because Google loves fresh, quality content (not in docs). **ANALYSIS** Google documents freshness systems only for queries that expect current information.
- **STAGE** Spot unwanted indexed pages; group an article's queries to fix bad cannibalization; start each experiment with a hypothesis and let the LLM report the test group's data.

STAGE For LLM traffic he trusts only the GA4 API with Looker Studio; GA4's LLM channels miss AI Overviews and AI Mode. **DOCS** Search Console counts that traffic as Web search. **STAGE** Tag Manager events can track each LLM (not in docs). Is good SEO good GEO? His agency's combined data says "it depends". **ANALYSIS** That mixed Google's AI features with other LLMs; compare each surface per page.

A second talk: when rankings rise and clicks don't

STAGE A business's rankings for valuable commercial keywords rose over a year while clicks stayed flat: "we rely on a metric that no longer stands up to scrutiny". AI takes part of the demand, informational queries first; in the speaker's brand data local terms now lose clicks too, and AI features spread to commercial results (not in docs).

DOCS Google wrote in August 2025 that total organic clicks were relatively stable, with traffic shifting between sites and some quick answers ending without a click. **ANALYSIS** No figures on either side: measure per site and query type.

STAGE Think like marketers before attribution: look at downstream impact on other channels. A cited Similarweb study: AI recommendations lead to later visits, direct, paid or otherwise (not in docs). **ANALYSIS** Scope: ChatGPT, US desktop, three industries, 2026; recommended brands 2.5 times likelier to get a visit within 7 days, 55.9% of it branded searches.

COMMUNITY CASE STUDY

A controlled test of AI visibility

STAGE A brand that had lost organic clicks worked only on its AI visibility for a set period: no campaigns, no budget changes. Inquiries rose significantly (no figure given); the brand had a record month. Its top-of-page share in paid search fell, so paid click-through rate should have dipped.

+36%

paid search click-through rate during the test

One unpublished test of one unnamed brand; no period or baseline given: an illustration, not a benchmark



A journey without attribution. The speaker's simplified path from an AI recommendation to a click; a brand that bids on its own name may see the brand search arrive as a paid click. A cited study found 82% of people click a brand they know, whatever its position (not in docs): Red C's 2018 eye-tracking study of shopping-type searches, so quote that scope.

What to report instead

STAGE "rankings and clicks don't really equal real outcomes anymore". Ask the paid search team for Google Ads click-through rate reports, and report conversion rates and traffic by channel. **DOCS** Google suggests tracking conversions and time on site, and judging the overall value of visits from Search, not clicks alone.

ANALYSIS The generative AI report shows no clicks: put organic next to direct traffic, branded paid search, conversions and revenue. Before your own AI visibility test, record period, baseline and other marketing changes; compare with a similar control area.

10


SECTION 10 · MASTERING THE MESSY MIDDLE · PABLO PÉREZ

How AI changes purchase decisions

Google's marketing research found that shoppers who use AI Mode or AI Overviews feel more confident: they hand AI the effort of a decision, not the choice, and still come back to Search and brand websites to double-check.

STAGE Google has used the "messy middle" model since 2019 for the stretch between a purchase's trigger and the purchase. People loop between exploring the category, its prices, brands and options, and evaluating them until confident enough to buy, torn between gathering more information and the effort of processing it.

DOCS Google's report *Marketing in the messy middle* calls exploration expansive and evaluation reductive; its study with The Behavioural Architects held in more than 20 countries.



Pinball versus funnel. Google's research as presented: without AI, shoppers bounce between the two mindsets; AI users mostly explore first, then evaluate, and double-check before risky purchases (not in docs). The speaker still called the funnel a common language for strategy, not how people decide.

The new research: shoppers "boosted" by AI

STAGE How AI changes decisions is the number one question Google's consumer insights team receives. The new study rests on conversations with consumers and, because people don't always do what they say, on behavioural panels run with research partners (not in docs). **PRESS** The DMEXCO 2026 programme (24 September) described it as tens of thousands of AI-moderated interviews and observed journeys across Europe.

31,000

in-market online shoppers in the original messy-middle study

UK, 2019; an earlier research wave

23,000

consumer conversations behind the AI findings

Presented October 2026; countries not stated

40,000

participants in behavioural panels, with research partners

One panel or the total: not made clear

STAGE The headline: people who use AI Mode or AI Overviews feel much more confident about purchase decisions. They hand AI the mental effort, not the choice: "that doesn't mean they are surrendering their choices. What's happening is that consumers are feeling boosted." Confused, overwhelmed shoppers come to feel like category experts (not in docs).

What AI gives shoppers (not in docs)

1 · EASE

STAGE Information condensed, choice overload made manageable; a photo instead of typing lowers the barrier.

2 · ASSISTANCE

STAGE A category cut to the features that matter (for a phone, battery and picture quality), plus what similar people chose.

3 · SUGGESTION

STAGE Ideas, brands and products the shopper hadn't thought of: curiosity costs less.

More searching, more brand websites

STAGE AI users run more searches and "also visit more websites of brands" (that part not in docs). Their journeys are sometimes longer but feel shorter, and they no longer feel stuck. In October 2026 they were a small but growing group (not in docs). **DOCS** Google says people use Search more than ever with the AI features (March 2025) and visit a greater diversity of websites with AI Overviews.

Reassurance: the double-check before buying

STAGE For high-risk or high-involvement purchases, ease, assistance and suggestion often aren't enough: after exploring with AI, people return to Google Search and the provider's own website to check trust in the company and details such as cancellation policies. A new or less familiar brand should expect that check and prepare for it (not in docs). **DOCS** Google's helpful-content questions ask about clear sourcing, expertise and background on the author or site.

SLIDE The brand heuristic: "A cognitive shortcut where consumers use a brand's name and reputation as a proxy for product quality, reliability, and value". **STAGE** In the reassurance step a familiar brand conveys a lot of confidence. **DOCS** The messy-middle report calls brand one of the strongest heuristics in consumer choice.

SLIDE Nine Google Trends panels (web search, 2016 to March 2026) showed trust checks rising in several languages, such as `is [brand] legit?`, `Ist [brand] seriös?` and `trustpilot`. **STAGE** The rise began around 2020 in the UK and most countries checked (not in docs): people taking the extra step.

What Google advised

- SLIDE** "Optimise for people to win in generative AI search" with unique, helpful, human-centric content: a human still decides. **DOCS** Google's guide to generative AI features says much the same.
- SLIDE** Partner with opinion formers: creators, experts, institutions (not in docs). **STAGE** Be present on Search and on your own site, where the brand speaks in its own voice.

RESEARCH ABOUT PEOPLE, NOT A RANKING SIGNAL

ANALYSIS Marketing research on how people decide, not on how Google ranks. Check the brand-visits finding against your own branded and direct traffic.

- Company details, reviews, policies and contact on indexable, linked pages.
- Track "[brand] legit" or "[brand] seriös" per market; answer on your own site.
- Name the features that decide a category; show what similar buyers chose.


11

SECTION 11 · HOW LONG DOES IT TAKE TO..? · GARY ILLYES

How long Google's processes take

Google shared estimates from its internal logs of how long its processes take, as information rather than targets. Finding and recrawling pages takes far longer than indexing them, and every step waits for the one before.

STAGE An experiment: estimates from internal logs, informative, nothing to obsess about. **SLIDE** Each chart showed a minimum, a typical value, an end point and dotted corner cases, with the footnote "Times expressed on the diagram are estimations based on internal analysis."



Fifteen processes on one time scale. Redrawn from Google's charts on a log scale, positions approximate: bar = minimum to usual end, dot = typical, dotted = corner cases up to "never". The meta-annotation minimum reads "200s" with no unit stated. Most of these figures are not in docs.

Crawling: finding and refreshing URLs

STAGE Google knows hundreds of trillions of URLs (October 2026) (not in docs) and crawls far fewer: a smaller crawl space is enough. URLs left unrefreshed for a very long time are dropped; news sites can be refreshed in seconds (not in docs). **ANALYSIS** Against the "hundreds of billions" of indexed pages in Google's ranking systems guide, about one known URL in a thousand is indexed, taken at face value.

STAGE On one audience member's very large site, the homepage was crawled about five times a day and soft-404 clusters every 160 to 190 days. **ANALYSIS** Link deep pages from strong hubs; check recrawl by depth in logs.

SLIDE A useful sitemap of a high-quality site is refetched within 14 days at most (not in docs); a lower-quality site's maybe never. **STAGE** robots.txt: a 24-hour objective, per RFC 9309; submitting it in Search Console forces a refresh. **DOCS** The same.

STAGE On 500 errors Google cuts crawl capacity within about four hours; increases take one to three weeks, and every site's capacity is recalculated within a month (that part not in docs). Buzz about a few URLs raises crawl demand for the whole site. **DOCS** Many 500, 503 or 429 replies lower the crawl rate until they stop. **ANALYSIS** A 5xx burst costs hours to cause and weeks to repair.

Rendering and indexing

STAGE Rendering happens right after crawling or later in a queue; JavaScript content reaches indexing in hours, at worst weeks; the queue clears within weeks (not in docs). Gary Illyes: "we keep saying that we render every single URL on the internet. I would say that that's not true" (not in docs).

ANALYSIS Google's JavaScript guide says every 200 page is queued for rendering; queued is not rendered. Put critical content, links and meta tags in the server HTML.

STAGE Meta annotations such as robots meta tags are a critical step: indexing waits for them (not in docs).

SLIDE "Indexing end to end" ends when a document reaches the serving index, tokenized (not in docs).

ANALYSIS Being found takes far longer than being indexed: for speed, link new pages from often-crawled ones and keep sitemaps accurate.

Changes you make

CHANGE	SAID ON STAGE	GOOGLE'S DOCS
404 or noindex	STAGE Out of the serving index in 1–3 weeks, often sooner.	
New canonical	STAGE Usually 1–3 weeks; can be seconds.	DOCS Clusters may hold up to two weeks. ANALYSIS The short end.
Site move	STAGE All signals recalculated: weeks (small site), 1–3 months on average, a year at worst.	DOCS Weeks for small sites; redirects at least a year.
Structured data	STAGE Hours to 1–2 weeks (not in docs); never if not trusted.	
Images, video	STAGE Images: hours to two days, under 150 s at best; video analysis: months (not in docs).	

STAGE Removals reach serving about every 2 hours; the image beside a result changes only after reindexing and an embeddings step (not in docs). Some manual actions are never lifted: nobody cleans up. **DOCS** Removals take up to a day and last six months; title links change after a recrawl (days to weeks); most reconsideration reviews take days or weeks.

WHAT TO DO WITH THESE NUMBERS

- ANALYSIS** Ship robots.txt changes a day ahead; for urgent ones, request a refresh.
- To hide a URL fast: the Removals tool (about 2 hours) plus a noindex or 404 (1–3 weeks).
- Judge a new title after days, a site move after one to three months; keep redirects a year.
- Manual actions lift within weeks of a real clean-up; core-update recovery takes months.


12

SECTION 12 · WRAPPING ALL UP · GARY ILLYES

AI in Search, used responsibly

The closing talk put AI back into proportion: machine learning has run in Search for decades, hallucinations can't be removed, and AI features use the same processes as classic results, so AI SEO is just SEO.

STAGE The closing message: use AI, but understand it well enough to use it responsibly. AI is an umbrella of technologies, one of which is machine learning, systems that learn from large amounts of data (not in docs). Machine learning is about 50 years old; Google has used it for "probably 30 years", since the statistical models behind "Did you mean" (not in docs). Some ranking features are plain machine learning models, which beat LLMs at some tasks.



AI is an umbrella. The closing talk's map of the terms, drawn from what was said on stage.

STAGE LLMs are deep learning on internet-scale data, mapping concepts by context in a vector space; predictive models guess the next word, generative ones produce output from a prompt (not in docs). BERT, long used in Search, is technically an LLM.

Hallucinations and deception

STAGE Any model can hallucinate, and current training can't stop it; grounding or RAG reduces hallucinations but cannot eliminate them (not in docs). Models, image models included, make things up when information is missing.

STAGE Any deception earns a page the lowest rater rating, and fake personas are a growing case: an author bio with an AI-generated headshot of a person who doesn't exist. Raters can't penalise a site, but their ratings become labels that improve the algorithms, so unchecked AI output harms indirectly. **DOCS** The rater guidelines (September 2025) rate made-up author profiles Lowest; no single rating changes a page's ranking directly. **ANALYSIS** Show only real authors with real photos and verifiable profiles, and check AI output before it goes live.

AI slop is scaled content abuse

STAGE Pages once churned out with Perl or PHP scripts are now churned out with LLMs, and cheaper tokens mean more AI slop. Google counts it as scaled content abuse and goes to great lengths to remove it. Citing its rater guidelines: "it's not the tool that was used that is the problem, but rather how the tool was used and what for." **DOCS** Google added this to its generative AI content guidance on 1 October 2026.

The closing slides

SLIDE Use AI to supercharge workflows, content, website and customer management, but responsibly: AI hallucinates, so don't add to the sea of AI slop.

"AI features on Google Search use exactly the same processes as traditional results."

Closing slide: "AI on Google is just SEO", no new acronym needed

STAGE AI Overviews and AI Mode sit on 25 to 30 years of Search infrastructure, with very few processes of their own. Gary Illyes: "To me, that means that SEO is not dead." Search changed every year, from the Florida update in 2003 to the Knowledge Graph (not in docs), each time with new opportunities; shiny things such as guessing fan-out queries distract from helpful content. Web Guide groups results into categories with summaries. From the community talks he judged that most SEOs have a very good handle on Search. **DOCS** A Search Labs experiment since July 2025, using a custom Gemini and query fan-out.

Three days, one pipeline

TOPIC	DAYS 1 AND 2	DAY 3
One index for AI	SLIDE Same crawl and index. Same snippets.	STAGE Very few processes of their own.
Query fan-out	DOCS No page per variant. Same index.	STAGE Like typed queries; don't overfocus.
Language	STAGE Country and language help ranking.	STAGE Detected first, matched at retrieval.
Markup and AI	STAGE The same data feeds AI features.	STAGE Not needed. ANALYSIS Its data can still help.
Rendering Tension	DOCS Every 200 page queued. Nearly all rendered.	STAGE Not every URL. ANALYSIS Queued ≠ rendered.
robots.txt	DOCS Cached for up to 24 hours.	SLIDE About 24 h. STAGE Per RFC 9309.
Server errors	DOCS 5xx slow crawling at once.	STAGE Capacity cut within about 4 hours.
AI against spam	STAGE More AI, highly accurate.	STAGE Faster, stronger spam updates.
MUM	DOCS Not used for general ranking.	ANALYSIS Listed on a slide: a system, not proof it ranks.
Measurement	SLIDE Business metrics. DOCS AI report: no clicks.	STAGE AI impressions reported; clicks aren't outcomes.
For people	SLIDE Optimise for people.	SLIDE The same, from marketing research.

13

SECTION 13 · TAKE IT BACK TO THE OFFICE

Day 3 action checklist

Checks that turn Day 3 into work, grouped for the content team, the development team and whoever reports on Search.

CONTENT TEAM

- ☐ Pages use the words and spellings the audience types; no synonym lists; variants checked in Search first
- ☐ Reviews and AI-assisted articles add own tests, photos and expertise beyond what an LLM would answer
- ☐ Product Q&As answer real customer questions from support tickets, reviews and sales chats
- ☐ Query groups are read as intent themes: one page per intent, not one per phrasing
- ☐ Company details, policies, reviews and contact are indexable pages; "[brand] legit" queries are tracked
- ☐ Only real authors with real photos and verifiable profiles; no invented personas or AI headshots

DEVELOPMENT TEAM

- ☐ Each page has one main language and a clear target country, with hreflang between the versions
- ☐ The back button works; Googlebot sees what users see; no doorway pages; the CMS is patched
- ☐ Articles allow `max-image-preview:large` with an image at least 1,200 px wide; WebSite markup names the site
- ☐ Feed and Product markup carry identical values; Storebot-Google is not blocked on shop paths
- ☐ Critical content, links and meta tags are in the HTML the server sends
- ☐ robots.txt changes ship a day early; redirects stay a year; urgent removals use the Removals tool

MEASUREMENT AND REPORTING

- ☐ The AI report is read as impressions; visits, conversions and revenue come from your own analytics
- ☐ Social accounts are platform properties; multimodal traffic is judged by page, country and device
- ☐ Large sites use the BigQuery bulk export; the branded split is spot-checked before LLM analysis
- ☐ After a core update: compare a full week after rollout with the week before, then review the whole site

14

SECTION 14 · BEYOND THE DOCUMENTATION

Said on stage, not in Google's docs

Points from Google's speakers that I could not find in Google's documentation. Not wrong for that, but not commitments either: treat them as informed hints and recheck before you build on them.

POINT	WHO	WHY IT MATTERS
STAGE Fan-out queries are not added to Search Console	John Mueller	ANALYSIS The docs fit: AI answers count under the user's own query.
STAGE Stop words are dropped; "siblings" such as Canon and Nikon are kept apart	John Mueller	[Canon camera] should not show Nikon cameras.
STAGE Quality is the most important signal for ordering candidates at retrieval	Gary Illyes	ANALYSIS The docs name quality for ranking, never for retrieval.
STAGE Search quality raters may also work for other search engines	Google	SLIDE Raters cannot affect any one site's rankings.
STAGE Scaled content abuse replaces link spam as the spam type worth talking about	Google	ANALYSIS Audit mass-produced pages before old link patterns.
STAGE The results page is an auction of result types, a metaphor, not ads	Gary Illyes	Text results remain the most common type.
STAGE An image used as a query is searched as embeddings	Gary Illyes	Day 2: Google also retrieves documents through embeddings.
STAGE A few flexible feed attributes instead of hundreds of narrow ones	Alex Jansen	Six were added for AI shopping.
STAGE The new attributes can also be read from schema.org markup	Alex Jansen	ANALYSIS Merchant Center maps only item_group_title; the feed is the documented route.
STAGE AI-powered configuration takes any language; more platforms are coming to platform properties	Ariel Kroszynski	Today: Instagram, TikTok, X and YouTube.
SLIDE On average, about 20 hours to discover a URL, 30 days to refresh it, 1.5 hours to index it	Gary Illyes	ANALYSIS Speed comes from discovery: links from often-crawled pages, accurate sitemaps.
STAGE Not every URL gets rendered	Gary Illyes	ANALYSIS Queued is not rendered: don't rely on rendering for critical content.
STAGE Grounding (RAG) reduces hallucinations but cannot remove them	Gary Illyes	No current training method removes them either.
STAGE Shoppers are "boosted" by AI: they hand over the effort, not the choice	Pablo Pérez	ANALYSIS Quote the 2026 research apart from the 2019 messy-middle study.

15

SECTION 15 · STAY IN TOUCH

Help complete this guide

This guide covers what I could capture in the room. If you heard something on Day 3 that isn't written here, tell me and I'll add it, so everyone who attended ends up with the full picture.

01 · ADD WHAT I MISSED

Heard something that isn't here?

These parts of Day 3 are the biggest gaps:

- Quality prevails, after lunch
- Uncovering Trustworthy Experiences on Discover: slides only, no recording
- Poster session M: Success Strategies That Work
- The Q&A panel and the closing

Each addition is labelled and checked against Google's documentation before it goes in, and credited to you by name if you like.

02 · THE FULL KNOWLEDGE BASE

Every Day 3 point, in the web edition

This guide is a selection. The web edition holds all of Day 3's points, about 700, with topic pages, session pages, a search and the verification status of every point. It is a folder that opens offline from `index.html`.

An agent-ready version of the same data lets AI tools answer questions about the event, with every answer traceable to its slide, session or document.

Message me and I'll share both.

LET'S KEEP IN CONTACT

Ibrahim Anjro

TEDX SPEAKER

WhatsApp **+49 1575 1702770**

LinkedIn [linkedin.com/in/ibrahem-angro](https://www.linkedin.com/in/ibrahem-angro)

Message me with additions, corrections, or to get the web edition.



WhatsApp



LinkedIn

REFERENCES

Sources

This guide is built from 33 photos of the slides and my recordings of 11 sessions, transcribed and checked against the slides where I had them. Every point labelled Docs was checked against these pages on 3 October 2026.

1. [Search generative AI control](#) — Google Search Console Help
2. [Optimizing your website for generative AI features on Google Search](#) — Google Search Central
3. [Introducing Search Generative AI performance reports in Search Console](#) — Search Central blog (3 June 2026)
4. [In-depth guide to how Google Search works](#) — Google Search Central
5. [A guide to Google Search ranking systems](#) — Google Search Central
6. [Creating helpful, reliable, people-first content](#) — Google Search Central
7. [List of Google's common crawlers](#) — Google
8. [How Google interprets the robots.txt specification](#) — Google
9. [Optimize your crawl budget](#) — Google
10. [How HTTP status codes affect Google's crawlers](#) — Google
11. [Search Central Deep Dive Europe 2026: Apparently we're going to Barcelona](#) — Search Central blog (6 July 2026)
12. [Build and submit a sitemap](#) — Google Search Central
13. [Robots meta tag, data-nosnippet, and X-Robots-Tag specifications](#) — Google Search Central
14. [Understand the JavaScript SEO basics](#) — Google Search Central
15. [Tell Google about localized versions of your page](#) — Google Search Central
16. [Control your snippets in search results](#) — Google Search Central
17. [AI features and your website](#) — Google Search Central
18. [Discover and your website](#) — Google Search Central
19. [Search Engine Optimization \(SEO\) Starter Guide](#) — Google Search Central
20. [Verify your site ownership](#) — Google Search Console Help
21. [Fix canonicalization issues](#) — Google Search Central
22. [How to move a site](#) — Google Search Central
23. [Managing multi-regional and multilingual sites](#) — Google Search Central
24. [General Guidelines](#) — Google Search Quality Rater Guidelines (PDF, 11 September 2025)
25. [Introduction to structured data markup in Google Search](#) — Google Search Central
26. [General structured data guidelines](#) — Google Search Central
27. [Structured data markup that Google Search supports](#) — Google Search Central
28. [Review snippet \(Review, AggregateRating\) structured data](#) — Google Search Central
29. [Merchant listing \(Product, Offer\) structured data](#) — Google Search Central
30. [Latest documentation updates](#) — Google Search Central
31. [Google Search's guidance on using generative AI content on your website](#) — Google Search Central
32. [Video SEO best practices](#) — Google Search Central
33. [Spam policies for Google web search](#) — Google Search Central
34. [How search results may differ based on accented characters and interface languages](#) — Search Central blog (1 September 2006)
35. [How Google knows what language to show in search results](#) — Google Search Help
36. [How we fought spam on Google Search in 2022](#) — Search Central blog (11 April 2023)

- [37. How we fought Search spam on Google in 2021](#) — Search Central blog (21 April 2022)
- [38. Google Search spam updates and your site](#) — Google Search Central
- [39. How Google Search organizes information](#) — Google Search (How Search Works)
- [40. Expanding AI Overviews and introducing AI Mode](#) — Google blog (5 March 2025)
- [41. Compare Trends search terms](#) — Google Trends Help
- [42. robots.txt report](#) — Google Search Console Help
- [43. Automatically generating and ranking results](#) — Google Search (How Search Works)
- [44. What are impressions, position, and clicks?](#) — Google Search Console Help
- [45. Grounding with Google Search](#) — Google AI for Developers (Gemini API docs)
- [46. Refine Google searches](#) — Google Search Help
- [47. Improving Search with rigorous testing](#) — Google Search (How Search Works)
- [48. Detecting spam to bring you relevant and reliable results](#) — Google Search (How Search Works)
- [49. Our approach to Search](#) — Google Search (How Search Works)
- [50. Google Search Status Dashboard](#) — Google Search Status Dashboard (Ranking incident history)
- [51. Google Search Status Dashboard](#) — Google Search Status Dashboard (September 2026 spam update)
- [52. Introducing a new spam policy for "back button hijacking"](#) — Search Central blog (13 April 2026)
- [53. Report spam, phishing, or malware](#) — Google Search Central
- [54. Google Search's guidance about AI-generated content](#) — Search Central blog (8 February 2023)
- [55. Google Search's core updates and your website](#) — Google Search Central
- [56. What site owners should know about Google's August 2019 core update](#) — Search Central blog (1 August 2019)
- [57. How we fought Search spam on Google in 2020](#) — Search Central blog (29 April 2021)
- [58. From the garage to the Googleplex](#) — Google
- [59. We knew the web was big...](#) — Official Google Blog (25 July 2008)
- [60. Introduction to Product structured data](#) — Google Search Central
- [61. Article \(Article, NewsArticle, BlogPosting\) structured data](#) — Google Search Central
- [62. Structured data for subscription and paywalled content \(CreativeWork\)](#) — Google Search Central
- [63. Provide a site name to Google Search](#) — Google Search Central
- [64. Visual Elements gallery of Google Search](#) — Google Search Central
- [65. Influencing your title links in search results](#) — Google Search Central
- [66. Help your readers find your site through preferred sources in Google Search](#) — Google Search Central
- [67. Universal search: The best answer is still the best answer](#) — Official Google Blog (16 May 2007)
- [68. Product data specification](#) — Google Merchant Center Help
- [69. Question and answer \[question_and_answer\]](#) — Google Merchant Center Help
- [70. Document link \[document_link\]](#) — Google Merchant Center Help
- [71. Related product \[related_product\]](#) — Google Merchant Center Help
- [72. Popularity rank \[popularity_rank\]](#) — Google Merchant Center Help
- [73. Item group title \[item_group_title\]](#) — Google Merchant Center Help
- [74. Variant option \[variant_option\]](#) — Google Merchant Center Help
- [75. Product highlight \[product_highlight\]](#) — Google Merchant Center Help
- [76. Product detail \[product_detail\]](#) — Google Merchant Center Help
- [77. About the Google StoreBot crawler](#) — Google Merchant Center Help
- [78. Allow Merchant Center to update product information automatically](#) — Google Merchant Center Help
- [79. Free listings for products](#) — Google Merchant Center Help

80. [New tech and tools for retailers to succeed in an agentic shopping era](#) — Google blog (11 January 2026)
81. [How we're helping retailers thrive with new Universal Commerce Protocol features and AI tools on Google](#) — Google blog (20 May 2026)
82. [Boost your holiday sales with these agentic commerce updates](#) — Google blog (16 September 2026)
83. [Introducing Query groups in Search Console Insights](#) — Search Central blog (27 October 2025)
84. [Streamline your Search Console analysis with the new AI-powered configuration](#) — Search Central blog (4 December 2025)
85. [Introducing social channels in Search Console](#) — Search Central blog (8 December 2025)
86. [See how content from social and video platforms performs on Google Search](#) — Search Central blog (7 July 2026)
87. [Platform properties roll out globally, plus a new social and video performance guide](#) — Search Central blog (29 July 2026)
88. [Analyze your social and video platform content performance in Search Console](#) — Google Search Central
89. [About platform properties in Search Console](#) — Google Search Console Help
90. [Announcing web multimodal Search performance reporting in Search Console](#) — Search Central blog (24 September 2026)
91. [Performance report \(Search results\): Overview and basic setup](#) — Google Search Console Help
92. [Performance report \(Search results\): Advanced filtering and comparison](#) — Google Search Console Help
93. [Generative AI performance report \(Search\)](#) — Google Search Console Help
94. [Britney Spears spelling correction](#) — Google (archived jobs page, 2011)
95. [Search with Google Lens in Chrome](#) — Google Chrome Help
96. [Search with an image on Google](#) — Google Search Help
97. [A deep dive into Search Console performance data filtering and limits](#) — Search Central blog (19 October 2022)
98. [Introducing the branded queries filter in Search Console](#) — Search Central blog (20 November 2025, updated 11 March 2026)
99. [About bulk data export of Search Console data to BigQuery](#) — Google Search Console Help
100. [Bulk data export: a new and powerful way to access your Search Console data](#) — Search Central blog (21 February 2023)
101. [Table guidelines and reference](#) — Google Search Console Help
102. [Query guidelines and sample queries](#) — Google Search Console Help
103. [AI in Search is driving more queries and higher quality clicks](#) — Google blog (6 August 2025)
104. [Top ways to ensure your content performs well in Google's AI experiences on Search](#) — Search Central blog (21 May 2025)
105. [AI-Recommended Brands Saw 2.5x More Site Visits: Similarweb](#) — Search Engine Journal (23 June 2026)
106. [82% of searchers choose a familiar brand for the first click \[study\]](#) — Econsultancy (29 October 2018)
107. [Marketing in the messy middle](#) — Think with Google (Google and The Behavioural Architects; Decoding Decisions, part 2)
108. [Beyond the Messy Middle: Mastering the consumer journey with AI](#) — DMEXCO 2026 event programme
109. [Update your robots.txt file](#) — Google Search Central
110. [Reduce the Google crawl rate](#) — Google
111. [Ask Google to recrawl your URLs](#) — Google Search Central
112. [Removals and SafeSearch reports tool](#) — Google Search Console Help
113. [Manual actions report](#) — Google Search Console Help
114. [Web Guide: An experimental AI-organized search results page](#) — Google blog (24 July 2025)
115. [Discover content policies](#) — Google Search Help

Search Central Live is a Google event. This is an independent attendee guide by Ibrahim Anjro, not an official Google publication. Statistics are as presented on stage, with their scope as stated.