

How Google Crawls in the Age of *AI Search*

A senior-level guide to Day 1 in Barcelona: what Google showed about AI Mode, how the crawler really works, why robots.txt breaks, and how crawl budget is decided. Every point labelled and checked against Google's documentation.

NOTES & GUIDE BY

Ibrahim Anjro

EVENT

Barcelona · 30 September 2026

EDITION

2 October 2026

Contents

—	Key takeaways	03
	The ten things to remember from Day 1	
01	Day 1 at a glance	04
	The agenda and what this guide covers	
02	The state of Search	05
	Google's AI Mode numbers and five ecosystem principles	
03	New publisher controls	07
	The generative AI opt-out, Search profiles, preferred sources	
04	How Search works, and where AI sits	09
	One pipeline for Search, AI Mode and Gemini; "Good SEO is good GEO"	
05	How crawling works	11
	Queue, scheduler, crawler, and one shared infrastructure	
06	Crawl errors	12
	DNS, firewalls, soft 404s, 5xx and 429, HTTP/3	
07	robots.txt, read the way Google reads it	14
	A puzzle solved URL by URL, and the group-closing trap	
08	Crawl budget	16
	Rate limit and demand, faceted URLs, and what each directive costs	
09	From the floor	18
	Duplicate content, pagination, images, accessibility	
10	Audit checklist	19
	Twelve checks to run on your own sites this week	
11	Help complete this guide	20
	Send what I missed, and get the agent-ready version	
—	Sources	21
	Google documentation used to verify this guide	

HOW TO READ THE LABELS

- SLIDE

 Shown on screen during the session
- STAGE

 Said on stage, from my session notes
- DOCS

 Google's published documentation
- PRESS

 A third-party report of what Google said
- ANALYSIS

 My interpretation, not Google's words

SUMMARY

Key takeaways

AI features run on the same crawl and index as classic Search. Whether Google reaches and fetches your pages still decides everything downstream, including AI answers.

01 Good SEO is good "GEO".

AI Overviews and AI Mode are built on the core ranking systems and ground answers in the Search index. There is no separate optimisation discipline to buy.

02 Expect traffic to move.

Google said outright that long-held traffic patterns are likely to fluctuate. Growth is in brainstorming, "which" and planning queries.

03 There is now a real AI opt-out.

A Search Console setting removes a property from AI Overviews, AI Mode and AI in Discover without, Google says, affecting regular rankings.

04 lms.txt does nothing for Google Search.

Google's own guide says such files "neither harm nor help". No special markup or content chunking is needed for AI features.

05 One crawling infrastructure serves Google's products.

Search, Ads, Shopping and Images share one crawling infrastructure, and one hostload limit per host.

06 Your server's health sets the crawl rate.

Connect time, time to first byte and 429 or 5xx responses decide how hard Google crawls. Firewalls that silently block Googlebot look like a failing server.

07 A crawler obeys exactly one robots.txt group.

Googlebot ignores the `*` group when a Googlebot group exists, the longest matching rule wins, and one stray line can merge or split groups.

08 Crawl demand is inherited by path.

When Google knows little about a URL, it uses the quality and popularity of its parent folder, then that folder's parent.

09 noindex costs crawl; disallow doesn't.

Google must fetch a page to read noindex. crawl-delay is ignored. Faceted URLs are the classic budget leak.

10 There is no duplicate content penalty.


Google picks one URL to keep. The real cost is losing control of which one, and crawling every copy.

01

SECTION 01 Day 1 at a glance

The day followed one arc: the state of Search, then the crawl pipeline, then the three ways sites waste or block crawling — errors, robots.txt and crawl budget.

11:00	- Welcome and opening keynotes - Google Search leadership	Slides captured
11:35	- What's new in the world of Search - Google	Slides captured
11:45	- How Search works and where's AI? - Google and Gary Illyes	Slides captured
13:10	- Lightning session A: Automation and AI - Community speakers	Not covered
14:05	- How crawling works - Google	Slides captured
14:35	- How crawling errors affect Search - Google	Notes only
15:10	- How Google interprets robots.txt - Google	Slides captured
15:30	- Lightning session B: Robots.txt - Community speakers	Not covered
15:40	- What would you do?: Crawling Edition - Google	Not covered
16:00	- How Google thinks about crawl budget - Cherry Pommawin, Search Advocate	Slides captured
16:20	- Lightning session C: Crawling - Community speakers	Not covered
16:35	- Q&A - Google panel	Not covered
17:25	- Daily wrap-up - Google	Not covered



The three-day arc. The event followed a page through Google Search, one stage per day, as shown in "How Search works".

02

SECTION 02 · KEYNOTE AND "WHAT'S NEW"

The state of Search

Google's message: AI is growing total search use rather than replacing it, and traffic will shift as a result.

"Search is never a solved problem. Old challenges evolve and new challenges are constantly popping up — because the internet and the world are always changing."

Elizabeth Reid, VP, Search — opening slide

SLIDE The theme guiding the future of Search: "**Making it effortless to find helpful information.**" Three forces were named as drivers: new platforms (closed-garden apps taking information needs away from the open web), new user preferences (authentic, visual, easy-to-scan content) and new technologies (AI).

The numbers Google showed

Google-reported figures from the slides. The slide footnotes cite Google internal data and Google-commissioned Ipsos global surveys; only the follow-up and brainstorming figures were marked as US. Treat them as Google's framing, not independent research.

5T+ searches on Google per year	2× / qtr AI Mode query volume, doubling every quarter since launch	~3× length of an AI Mode query vs. a traditional one	1 in 6 AI Mode searches are multimodal (voice or image)
+40% follow-up queries in AI Mode, per month on average US, since launch	1.5× faster volume growth for queries of 5+ words Global English queries	94% of frequent LLM users are also frequent Google users	71% of shoppers on Google are open to new brands Ipsos survey

Four new ways people search in AI Mode

EXPLORE Brainstorming queries growing 30% faster than AI Mode overall in the US since launch. "Where to", "ideas for", "where should I".	LEARN Study guides and quizzes, topic deep dives, professional development.	DECIDE Queries starting with "which" growing 40% faster than AI Mode queries overall over the last six months.	DO Planning queries — restaurants, itineraries, finances, fitness — growing 80% faster than AI Mode queries overall over the last six months.
---	---	---	--

SLIDE Google also showed **Search Agents** in AI Mode: standing requests such as "tell me if Samsung Frame 2000 goes on sale" or "alert me when the government shutdown ends".

Google's five principles for the ecosystem

#	PRINCIPLE	WHAT GOOGLE SAID SLIDE
1	Consistent testing	Features are tested, then woven into core Search when ready.
2	The SERP will evolve	Beyond organic links and ads: videos, cards and more.
3	Multiple user needs	Some want links, some quick answers, some depth. Google says it will keep prioritising traffic to the ecosystem.
4	Incentivise quality	"Made for Search" content will not be successful.
5	Traffic patterns may fluctuate	Long-held patterns are likely to change, "creating new opportunities for all sites".

ANALYSIS

Principle 5 is the line to plan around: Google told a room of SEOs to expect traffic to move. The growth is in brainstorming, "which" and planning queries — mid-funnel, comparison-style needs that a page answering one keyword fits poorly. Build pages that help someone decide, not just pages that match a term.

03

SECTION 03 · WHAT'S NEW

New publisher controls

Google listed eight ways AI Search sends people to websites, and separately showed a new opt-out control for generative AI. The opt-out is the one most site owners can act on today.

SLIDE The full list on screen: preferred sources expansion, fresh perspectives and prominent links, subscription content, in-line quotes in AI Mode and AI Overviews, rich linking and embedded web, context on content sources, "Highly Cited" labels, and Search profiles.

The generative AI opt-out in Search Console

SLIDE "Opt-out control for generative AI: direct control over AI Overviews and AI Mode without impacting web rankings." The setting is called **Search generative AI** and has three states.

Search generative AI

Property setting · Search Console

☐ **Inherit from parent property**
Default. Follows the closest parent that changed it

☒ **Include**
Links and content can appear and ground AI answers

☐ **Exclude**
No links, no grounding, no AI impressions or traffic

IT CONTROLS

- ✓ AI Overviews
- ✓ AI Mode
- ✓ Generative AI features in Discover

IT DOES NOT TOUCH

- ✗ Regular rankings or inclusion
- ✗ Gemini training (use Google-Extended)
- ✗ Individual pages (property level only)

What the switch does. Illustration based on the slide and Google's Search Console help page; not a screenshot of the product.

DOCS Google's help page adds the operational details. The control "isn't used as a ranking or inclusion signal" for regular results. Changes generally apply within 1–2 days, though some content takes longer. A child property inherits from its closest parent that changed the setting, unless you override it.

How to decide **ANALYSIS**

YOUR SITUATION	LIKELY CHOICE	WHY
Lead generation, local services, e-commerce	Include	AI answers cite and link sources; excluding removes you from a growing surface.
Paywalled or licensed content you sell	Test exclude on a child property	Protects the asset, but you lose AI citations. Measure before rolling out.
You only want to stop Gemini training	Include + Google-Extended	The opt-out does not control training.
You want out of AI and regular snippets	nosnippet	Wider effect: it also removes normal snippets.

Before changing anything, read the **Generative AI performance report** in Search Console. **DOCS** Google's launch post (3 June 2026) lists impressions, pages, countries, devices (Search only) and dates, and no click or query metrics — it measures exposure, not value.

Search profiles, preferred sources and subscriptions

SEARCH PROFILES

SLIDE A claimable profile: photo, bio, social accounts, pinned posts, links. **DOCS** Launched 4 June 2026, US first. Since 16 September 2026, US publishers and creators qualify with 10,000 followers across YouTube, Instagram, X or TikTok.

PREFERRED SOURCES

SLIDE Sites users mark as preferred get more visibility in Top Stories and a label in AI Mode and AI Overviews.

LINKED SUBSCRIPTIONS

SLIDE Users link subscriptions so answers highlight paywalled content they already pay for. One publisher saw 34% more engagement.

ANALYSIS


All three depend on users acting on your behalf. That rewards brands with a loyal audience and does little for anonymous lead-generation sites. Brand building is now ranking-adjacent work, not only marketing.

04

SECTION 04 · HOW SEARCH WORKS AND WHERE'S AI?


How Search works, and where AI sits

"If you know how Search works, you largely also know how AI features work." AI Mode and AI Overviews use the same crawl and index as classic Search. They differ only at serving.



One pipeline, three products. Redrawn from the session's matrix slide, which also served as the roadmap for the three days.

The AI already inside Search is not new



SLIDE Ranking signals differ by result type: **web pages** (text, links, passages), **images** (resolution, colour, associated text), **news** (freshness, originality, diversity), **local** (location, type, rating, reviews, hours) and **video** (language, text from speech).

"Good SEO is good GEO"

STAGE Gary Illyes argued that "GEO" is a new label invented to create a new field. The work doesn't need it: understand how SEO works and you are covered.

SLIDE "Don't worry about what to call it. Good SEO is good GEO, AEO." Three reasons: generative AI features are built on the core ranking systems; query fan-out expands the original query to find related information; and generative AI features highlight content indexed by Google Search.

DOCS Google's optimisation guide for generative AI (updated 10 July 2026) agrees: optimising for generative AI search "is optimizing for the search experience, and thus still SEO". It describes query fan-out as running several related searches at once — a question about lawn weeds also searches herbicides and prevention.

Four myths Google addressed

MYTH	GOOGLE'S ANSWER SLIDE
Optimise for AI over readers	Optimise for people. No need to obsess over keywords or AI phrasing, no need to chop content, "LLMS.txt isn't necessary".
Build content for every possible need	Prioritise unique perspectives and expertise beyond common knowledge.
Your website is no longer relevant	Keep content accessible: crawlable, well structured, fast and easy to read, for readers and AI tools.
Old metrics don't work in the AI era	Focus on real business goals and measure what matters to your business.

AUDIENCE QUESTION · "CAN GOOGLE DISTINGUISH BETWEEN AI AND HUMAN WRITTEN CONTENT?"


SLIDE "ML based ranking algorithms and signals are trained on content by humans for humans. They 'understand' and promote natural content better."

ANALYSIS This is not a claim that Google detects AI text. It says ranking models favour text that reads as natural to people. Google's policy is consistent: AI-assisted content is allowed, while mass-produced pages made mainly to rank fall under its scaled content abuse policy. The risk is in scale without value, not in the tool.

05

SECTION 05 How crawling works

Google runs one crawling system for all its products, and its first job is not to harm your server. That is why speed and errors on your side decide how much Google crawls.



Google's crawl pipeline. Redrawn from the "How crawling works" and crawl budget slides.

- **SLIDE** The crawler has "multiple tasks": fetch from the internet, "ensure we don't break the internet", and enforce robots.txt policies. Fetched data goes on to indexing.
- **SLIDE** The scheduler is "shared scheduler infrastructure". It decides what to fetch and when; each team sets the scheduling parameters for its own user agent.
- **SLIDE** Search, Ads, Shopping and Images all request through one "centralized crawling infrastructure". Its primary mandate: fetch "while strictly preventing the overloading of external servers".
- **SLIDE** Discovery runs through links: a news homepage links to `/sport` and `/tech`, which link to their articles.

ANALYSIS · TWO CONSEQUENCES

Orphan pages are slow pages. A URL with no internal links depends on sitemaps alone to be found, so discovery is slow and demand stays low.

Every Google fetch counts against the same host. Ads and Shopping fetches share host capacity with Googlebot. A merchant feed or ad landing pages hitting a slow server can reduce how much is crawled for Search.


06

SECTION 06 · HOW CRAWLING ERRORS AFFECT SEARCH

Crawl errors: what Google sees when it can't reach you

The most damaging crawl errors are the ones site owners never see. A firewall that blocks Googlebot shows up on Google's side only as failed connections — and Google reacts by crawling less.

STAGE DNS problems and firewalls are blind spots: Google doesn't know a firewall is the cause, it only sees requests fail. Soft 404s were also raised.



Where a fetch can fail. Each layer between Googlebot and your page can fail in a way that looks like a sick server.

How Google reacts to each response **DOCS**

WHAT GOOGLE GETS	TREATED AS	EFFECT
5xx server error	Overload signal	Crawling slows; indexed URLs are kept for a while, "but eventually dropped".
429 Too Many Requests	Overload signal	Crawling slows, like 5xx .
4xx other than 429	Normal "not found"	"No effect on crawl rate"; the URL leaves the index.
200 with empty or error page	Soft 404	Excluded from the index — and still crawled, wasting budget.
Timeout, reset, DNS failure	Network / DNS error	Treated like 5xx: crawling slows at once, and indexed URLs that stay unreachable drop out of the index within days.
robots.txt returns 5xx	Rules unreadable	Crawling stops for 12 hours, then the cached copy is used for up to 30 days.

HTTP/3

STAGE From my notes, session not recorded: Googlebot does not support HTTP/3 today. **DOCS** Consistent: Google's crawlers use HTTP/1.1 and HTTP/2, pick whichever crawls best, and may switch between sessions. HTTP/2 can save server resources but gives no ranking benefit.

ANALYSIS

This doesn't mean "don't use HTTP/3" — browsers benefit from it, and servers offer it alongside HTTP/2 and HTTP/1.1. The rule is: never run a host that only works over HTTP/3, and make sure CDN or firewall rules written for HTTP/3 traffic don't break the protocols Googlebot uses.

How to find these errors

1. Search Console → **Settings** → **Crawl stats** → **Host status**: robots.txt fetch, DNS resolution and server connectivity. Any red in the last 90 days deserves a ticket.
2. In the same report, break requests down **by response**. A rising share of `5xx` or `429` means Google is being told to slow down.
3. Ask whoever runs the CDN or WAF for blocks on **Google's published crawler IP ranges**. Verify Googlebot by reverse DNS, never by user agent alone.
4. For temporary blocks, return `503` or `429` — never a `200` page with a captcha or error message, which becomes a soft 404 and can replace your real content in the index.
5. Review **Soft 404** in the Pages report. On local and service sites, thin location pages with little unique content are the usual cause.

07

SECTION 07 · HOW GOOGLE INTERPRETS ROBOTS.TXT

robots.txt, read the way Google reads it

Most robots.txt mistakes come from three rules people get wrong: a crawler obeys only one group, the longest matching path wins, and `$` anchors the whole URL — parameters included.

The puzzle from the slide

```
# Let's craft a robots.txt!

user-agent: *
disallow: /
allow: /$

user-agent: bingbot
user-agent: googlebot
disallow: /
allow: /$
allow: /politics/
allow: /*/live/

user-agent: google-extended
disallow: /subscriptions/
allow: /
```

DOCS The rules that decide it:

- 1. One group per crawler.** Several `user-agent` lines can share a group. A crawler follows only the most specific group naming it — so Googlebot ignores the `*` group here entirely.
- 2. Longest path wins.** Within the group, the rule with the longest matching path applies.
- 3. Ties go to allow.** If an allow and a disallow match equally, the less restrictive wins.
- 4. `*` and `$`.** `*` matches any characters; `$` marks the end of the URL.

What Googlebot may crawl under this file **ANALYSIS**

URL	DECIDING RULE	GOOGLEBOT
<code>/</code>	<code>allow: /\$</code>	Allowed
<code>/?utm_source=newsletter</code>	<code>disallow: /</code> — <code>/\$</code> needs the URL to end right after <code>/</code>	Blocked
<code>/index.html</code>	<code>disallow: /</code>	Blocked
<code>/politics/eu-vote</code>	<code>allow: /politics/</code> (longer than <code>/</code>)	Allowed
<code>/politics</code>	<code>disallow: /</code> — no trailing slash, no match	Blocked
<code>/sports/live/</code>	<code>allow: /*/live/</code>	Allowed
<code>/live/</code>	<code>disallow: /</code> — pattern needs a segment before <code>/live/</code>	Blocked
<code>/sports/live-score</code>	<code>disallow: /</code> — no <code>/live/</code> with a slash	Blocked

Every other crawler falls into the `*` group and may fetch only the bare homepage. `google-extended` is not a crawler — see below.

The group-closing trap

SLIDE A later slide of the same talk: "Different crawlers behave differently! ... Unpredictability is never a good time."

HOW GOOGLEBOT READS IT

```
User-agent: bingbot
Content-signal: all-access=yes
User-agent: googlebot
Disallow: /
```

Both bingbot and googlebot blocked

Unknown line ignored; one group, one rule

HOW BINGBOT READS IT

```
User-agent: bingbot
Content-signal: all-access=yes
User-agent: googlebot
Disallow: /
```

bingbot not blocked

Group closed early; bingbot ends up with no rules

Same file, two readings. Bars show how each crawler groups the lines, per the talk's slide.

ANALYSIS

`Content-Signal` lines are now common because a large CDN provider adds them to its managed robots.txt files. Google supports only `user-agent`, `allow`, `disallow` and `sitemap` and ignores the rest. Never place any other directive between `user-agent` lines: put it after a group's rules, and test the file with each search engine's tools (Bing's robots.txt tester, Google's open-source robots.txt parser).

Google-Extended, crawl-delay and llms.txt

Google-Extended

DOCS A control token, not a crawler. Decides whether content may train future Gemini models and ground Gemini apps and Vertex AI. It "does not impact a site's inclusion in Google Search nor is it used as a ranking signal".

crawl-delay

DOCS Not supported by Googlebot; silently ignored. robots.txt is cached for up to 24 hours, and only the first 500 KiB is read.

llms.txt

SLIDE "LLMS.txt isn't necessary." **STAGE** Speakers at the event were openly dismissive. **DOCS** Such files "neither harm nor help" — Google Search ignores them.


08

SECTION 08 · CHERRY PROMMAWIN, SEARCH ADVOCATE

Crawl budget

Crawl budget is two separate limits: how much your server can take, and how much Google wants. Speed raises the first; quality and popularity raise the second. Most sites only ever hit the second.


STAGE In one line: crawl budget is the attention span Google gives a website — and better performance increases it.



The formula. From the session slides: "Crawl Budget = Crawl Rate Limit & Crawl Demand".

Quality is inherited by path

SLIDE If quality or popularity is unknown, the parent root's aggregate quality or popularity is used, then that path's parent's, and so on.



Illustrative example. How path inheritance could play out for new content, based on the slide's rule. Paths are hypothetical.

ANALYSIS

Put new high-value content where Google already trusts you. A new folder under a section Google rates poorly starts with that section's low demand.

Who needs to care

DOCS Google's crawl budget guide is written for sites with **1 million+** unique pages changing weekly, **10,000+** pages changing daily, or many URLs stuck in "Discovered – currently not indexed". **ANALYSIS** On a smaller site, slow indexing is almost always a demand problem (quality), not a capacity problem.

What burns it: the six faceted-URL traps

SLIDE Server errors unrelated to load, useless pages and resources, and infinite URL spaces. Faceted navigation was the worked example.

TRAP	EXAMPLE FROM THE SLIDE	FIX
1 · Same facets, different order	<code>/shoes/color-red/size-10</code> <code>/shoes/size-10/color-red</code>	Always build facet URLs in one fixed order
2 · Impossible combinations	<code>/laptops/brand-apple/os-windows</code>	Return <code>404</code> when a combination has no results
3 · Too many facets	<code>/cameras/type-dslr/brand-canon/...</code>	Make one or two facet levels crawlable; block deeper ones
4 · Facets on pagination	<code>/phones/.../page-2/sort-price-asc</code>	Fixed parameter order; block sort parameters
5 · Facets + internal search	<code>/search?q=blender&brand-ninja&...</code>	Disallow internal search URLs in robots.txt
6 · Default-value facets	<code>/televisions/size-55inch/resolution-4k</code>	Drop default values; link only to the normal form

DOCS Google's faceted navigation guide prefers prevention: block filter URLs in robots.txt and keep crawlable only item pages plus one unfiltered listing page — or use URL fragments, which Google generally doesn't crawl. `rel="canonical"` and `nofollow` are weaker, slower options.

How to manage it

01

Use HTTP caching.
Serve `304` for unchanged pages; ETag preferred.

02

Have good navigation.
Clear internal links for discovery and priority.

03

Restrict faceted and action URLs — add to cart, wishlist, print, sort.

04

Improve or remove useless content.
Return `404` / `410` for removed pages.

What each directive costs

DIRECTIVE	CRAWL BUDGET SLIDE	THE CATCH DOCS ANALYSIS
disallow	Not consumed — never fetched	Can still be indexed without content if linked, and Google can't see its <code>noindex</code> . Not for temporarily shifting budget.
noindex	Consumed — page must be fetched	The right tool to keep a page out of the index; the wrong tool to save crawl.
nofollow	Can be consumed	A hint, not a block. URLs found through other links are still crawled.
crawl-delay	Consumed — ignored	No throttling effect. To slow Google temporarily, return <code>503</code> or <code>429</code> .

09

SECTION 09 · FROM MY NOTES

From the floor

Points said on stage that my notes do not tie to a session, checked against Google's documentation where possible.

No duplicate content penalty

STAGE Gary Illyes: there is no such thing as a duplicate content penalty. **PRESS** John Mueller, quoted by Search Engine Journal in April 2026: "no penalty or ranking demotion if you have multiple URLs going to the same content". **ANALYSIS** The real costs: Google may choose a canonical you didn't want, and every copy is crawled. Scraped content is a separate spam issue.

Pagination and "load more"

STAGE Adopt load-more patterns carefully; they're complex and mistakes are costly — check your logs. **DOCS** Crawlers "don't 'click' buttons". Each page needs its own URL and a plain `<a href>` to the next; don't canonicalise to page 1; `rel=next/prev` is no longer used.

Images matter

STAGE Images are important. **SLIDE** One in six AI Mode searches is multimodal, and a slide listed resolution, colour and associated text as image ranking signals. **ANALYSIS** Original photos, descriptive file names, alt text and captions now feed AI answers too.

Accessibility

STAGE Check ARIA and accessibility. **ANALYSIS** Google hasn't called ARIA a ranking signal. The case is that AI agents and assistive tools read pages through the same structure: real buttons, labelled forms, semantic headings.

Getting Google to build something

STAGE To get a feature added — an addition to an API, for example — request it publicly and in volume. Broad demand moves the roadmap.

SEO is more than content

STAGE SEO has many parts beyond content. **SLIDE** "Optimizing for people is optimizing for Generative AI Search."

Is SEO dead? Google's answer on stage was no: AI features run on the same crawl, index and ranking systems.

10

SECTION 10 · TAKE IT BACK TO THE OFFICE

Day 1 audit checklist

Twelve checks that turn Day 1 into work. Most take minutes in Search Console and one server log file.

ACCESS AND ERRORS

- ☐ Crawl stats → Host status shows no DNS, robots.txt or connectivity failures in the last 90 days
- ☐ Crawl stats by response: the share of `5xx` and `429` is flat or falling
- ☐ CDN / WAF rules don't block Google's published crawler IP ranges; Googlebot verified by reverse DNS
- ☐ Maintenance and temporary blocks return `503`, not a `200` error page
- ☐ Every "Soft 404" in the Pages report is improved, merged, or returned as `404` / `410`
- ☐ Every host answers over HTTP/1.1 and HTTP/2, whatever the HTTP/3 setup for browsers

ROBOTS.TXT

- ☐ Each crawler you care about gets the rules you expect — remembering it obeys only its most specific group
- ☐ No unknown lines (`Content-Signal` , `crawl-delay`) sit between `user-agent` lines
- ☐ The homepage with tracking parameters isn't accidentally blocked by a `/$` rule
- ☐ `Google-Extended` is set deliberately, knowing it governs Gemini, not Search

CRAWL EFFICIENCY AND AI EXPOSURE

- ☐ Filter, sort, internal search and action URLs are blocked or normalised; logs show Googlebot spending most requests on pages you want indexed
- ☐ The Generative AI performance report is reviewed before any change to the Search generative AI setting

11

SECTION 11 · STAY IN TOUCH

Help complete this guide

This guide covers what I could capture in the room. If you heard something on Day 1 that isn't written here, tell me and I'll add it, so everyone who attended ends up with the full picture.

01 · ADD WHAT I MISSED**Heard something that isn't here?**

These sessions are the biggest gaps:

- Lightning session A: Automation and AI
- Lightning session B: Robots.txt
- What would you do?: Crawling Edition
- Lightning session C: Crawling
- The Q&A

Send me the point, who said it, and a photo of the slide if you have one. Corrections are just as welcome.

Each addition is labelled and checked against Google's documentation before it goes in, and credited to you by name if you like.

02 · COMING NEXT**An agent-ready version**

I'm creating a version of these guides that AI agents can consume directly. If you use agentic tools, you'll be able to load it as a knowledge base and ask it questions about the event.

Every point will keep its source label and link, so answers can be traced back to the slide, the session or Google's documentation.

Get in touch and I'll share the files and access to the tool as soon as they're ready.

LET'S KEEP IN CONTACT

Ibrahim Anjro

TEDX SPEAKER

WhatsApp **+49 1575 1702770**

LinkedIn [linkedin.com/in/ibrahem-angro](https://www.linkedin.com/in/ibrahem-angro)

Message me with additions, corrections, or to get the agent-ready version.



WhatsApp



LinkedIn

REFERENCES

Sources

This guide is built from 45 photos of the slides, the printed agenda and my notes from the room. Every point labelled Docs was checked against these pages on 2 October 2026.

1. [Search generative AI control](#) — Google Search Console Help
2. [Optimizing your website for generative AI features on Google Search](#) — Google Search Central
3. [Introducing Search Generative AI performance reports in Search Console](#) — Search Central blog (3 June 2026)
4. [A new profile to help publishers and creators highlight their work on Search](#) — Google blog (4 June 2026)
5. [3 new ways we're improving Search profiles for publishers](#) — Google blog (16 September 2026)
6. [Manage your Search profile](#) — Google Search Help
7. [In-depth guide to how Google Search works](#) — Google Search Central
8. [A guide to Google Search ranking systems](#) — Google Search Central
9. [MUM: A new AI milestone for understanding information](#) — Google blog (18 May 2021)
10. [Creating helpful, reliable, people-first content](#) — Google Search Central
11. [Crawling infrastructure](#) — Google
12. [Overview of Google crawlers and fetchers \(user agents\)](#) — Google
13. [List of Google's common crawlers](#) — Google
14. [How Google interprets the robots.txt specification](#) — Google
15. [Optimize your crawl budget](#) — Google
16. [Managing crawling of faceted navigation URLs](#) — Google
17. [How HTTP status codes affect Google's crawlers](#) — Google
18. [Debug network and DNS errors for Google's crawlers](#) — Google
19. [Troubleshoot Google Search crawling errors](#) — Google Search Central
20. [Qualify your outbound links to Google](#) — Google Search Central
21. [Evolving "nofollow" – new ways to identify the nature of links](#) — Search Central blog (10 September 2019)
22. [Pagination, incremental page loading, and their impact on Google Search](#) — Google Search Central
23. [What is canonicalization](#) — Google Search Central
24. [Demystifying the "duplicate content penalty"](#) — Search Central blog (12 September 2008)
25. [Crawling December: HTTP caching](#) — Search Central blog (9 December 2024)
26. [Crawling December: CDNs and crawling](#) — Search Central blog (24 December 2024)
27. [Google Says It Can Handle Multiple URLs To The Same Content](#) — Search Engine Journal (8 April 2026)
28. [Search Central Deep Dive Europe 2026: Apparently we're going to Barcelona](#) — Search Central blog (6 July 2026)
29. [Inside Googlebot: demystifying crawling, fetching, and the bytes we process](#) — Search Central blog (31 March 2026)
30. [AI, personalization and the future of shopping](#) — Google blog (3 March 2025)
31. [Verify requests from Google crawlers and fetchers](#) — Google
32. [robots.txt report](#) — Google Search Console Help
33. [Minimize A/B testing impact in Google Search](#) — Google Search Central

Search Central Live is a Google event. This is an independent attendee guide by Ibrahim Anjro, not an official Google publication. Statistics are as presented by Google on stage.