

SEARCH CENTRAL LIVE · DEEP DIVE EUROPE 2026

DAY 2 FIELD GUIDE

How Google Indexes in the Age of **AI SEARCH**

A senior-level guide to Day 2 in Barcelona: how Google renders and deduplicates pages, which robots rules steer AI answers, and what gets indexed, plus the community talks on site moves and media. Every point labelled and checked against Google's documentation.



NOTES & GUIDE BY

Ibrahim Anjro

Barcelona · 1 October 2026

Deep Dive edition · an independent attendee guide

DIVE LOG • DAY 2 FIELD GUIDE

Contents

—	Key takeaways	03
	The fourteen things to remember from Day 2	
01	Day 2 at a glance	04
	The agenda, a map of the day and the three-day arc	
02	From HTML to extracted elements	06
	Parsing, link extraction, four questions from the welcome Q&A	
03	Controlling indexing	08
	Every robots rule and snippet control, what each does to AI answers	
04	Rendering and JavaScript	10
	The render loop, four common problems, three community talks	
05	Understanding what's on a page	13
	Main content, tokens, the chunking myth, soft 404 detection	
06	Duplicates and canonicals	15
	Clusters, canonical selection, canonical graphs, two site moves	
07	Feature extraction: structured data, images, video	19
	Why markup matters, images, video, AI images, two media talks	
08	Language, country and hreflang	22
	hreflang code mistakes, localisation, non-Latin scripts, posters	
09	Signals and index selection	25
	What is calculated at indexing; the Page indexing report	
10	Inside the index	27
	Tokens, posting lists, embeddings; one index and one set of snippets for AI	
11	Google Trends for SEO	28
	Share, not volume; three uses; what is new	
12	What changed since Day 1	29
	Two contradictions, and what Day 2 added to Day 1	
13	Day 2 action checklist	30
	Seventeen checks to run this week	
14	Said on stage, not in Google's docs	31
	Points to treat as informed hints	
15	Help complete this guide	32
	Send what I missed; get the web edition	
—	Sources	33
	Google documentation and press reports used to check this guide	

HOW TO READ THE LABELS

SLIDE Shown on screen during the session

DOCS Google's published documentation

ANALYSIS My interpretation, not Google's words

STAGE Said on stage, from my recordings or a second attendee's

PRESS A third-party report of what Google said

(*not in docs*) Said at the event; not in Google's docs

SUMMARY

Key takeaways

Day 2 followed a fetched page into the index. Most of what decides whether it is indexed, and how it is shown, in AI answers too, happens there.

- 01 Rendering is a separate, expensive pass.**
Every 200 page without noindex is queued for a detached render. Put what indexing needs in the server HTML.
- 02 Googlebot neither clicks nor scrolls.**
Content behind a click or a scroll listener never loads. The viewport is about 10,000 pixels tall.
- 03 A robots.txt block is not a noindex.**
noindex works only on a page Google may crawl, and removing it with JavaScript doesn't work.
- 04 Snippet rules are now AI rules.**
nosnippet and max-snippet also limit AI Overviews and AI Mode, which are built from snippets.
- 05 Where text sits matters; chunking doesn't.**
Footer words weigh less than main content. Gemini reads millions of tokens; 100-word chunks miss the point.
- 06 Duplicates become one representative plus alternate names.**
One URL per cluster is indexed and gets the others' signals; old domains in site: queries are alternate names.
- 07 Google decides the canonical.**
rel=canonical is often wrong, so it is one input. When your signals disagree, Google can't tell what you want.
- 08 Structured data is still the precise, cheap route.**
Google prefers markup to LLM extraction, and the same data feeds AI features. There is no special AI markup.
- 09 hreflang fails on small mistakes.**
Every page needs a self-reference, return links and ISO codes: sv not se, GB not UK, never EU.
- 10 Index selection is a quality gate.**
'Crawled – currently not indexed' is mostly about quality; 'Discovered' means Google doesn't want to crawl it yet.
- 11 There is no separate AI index.**
AI Overviews and AI Mode query the same index and use the same token-based snippets as web results.
- 12 A video needs a watch page and a visible player.**
Without a watch page, or below the fold, it isn't indexed, Gary Illyes said. The docs want the player at load.
- 13 Consolidate by intent; never redirect to the homepage.**
Merge pages that serve one intent and redirect only to a same-intent page. Removed URLs get a 404 or 410.
- 14 Google Trends shows share, not volume.**
A line falling since 2004 can hide rising searches. Compare terms in one chart; take volumes elsewhere.

01

SECTION 01


Day 2 at a glance

Day 2 followed a page from the crawler's hand-over into the index, then looked inside the index and at Google Trends. Community lightning talks ran in between. This guide selects from nearly 1,000 points; the web edition holds every one.

10:15	 Welcome to indexing day! Gary Illyes and Cherry Prommawin, Search Relations	Transcript · Slides
10:25	 How is HTML interpreted Cherry Prommawin, Search Relations	Transcript · Slides
10:30	 Controlling indexing John Mueller, Search Relations Team Lead	Transcript · Slides · Video
10:40	 Lightning session D: Rendering and JavaScript Erin Sparling, Sören Bendig, Rebecca Yu and Natalia Venditto	Transcript · Slides
11:15	 What is Google friendly JavaScript Erin Sparling, Partner Technology Manager	Transcript · Slides · Video
11:30	 Understanding what's on a page Gary Illyes, Search Relations	Transcript · Slides · Video
11:55	 Handling web duplication John Mueller, Search Relations Team Lead	Transcript · Slides
12:05	 Lightning session E: Managing Duplicates and Site Moves Tobias Schwarz, Martyna Aġanoġlu and David Carrasco Pamies	Transcript · One slide
13:30	 Finding the gold nuggets: structured data, media, and more! Gary Illyes, Search Relations	Transcript · One slide
13:35	 What is Structured Data and why we need it on the internet. Ryan Levering, Software Engineer	Transcript · Slides
13:50	 Using images to your advantage and Engaging Search users with videos Gary Illyes, Search Relations	Transcript · One slide
14:00	 Lightning session F: Media a community speaker and Patrick Domanico	Transcript
14:15	 Focusing on Internationalisation and Localisation Google	Transcript · Slides
14:30	 Lightning session G: Internationalisation Alizée Baudéz and a second community speaker	Transcript
14:45	 Poster session H: Indexing Community speakers	Transcript
15:30	 Calculating (some) signals Google	Transcript
15:40	 Deciding what goes in the index? Google	Transcript
15:50	 Lightning session J: Indexing and Troubleshooting Community speakers	Not covered
16:00	 What would you do?: Indexing Edition Google	Not covered

16:15	☒ How does the index look like? Google	Transcript
16:20	☒ Google Trends Omri Weisman, Engineering Manager, Google Trends	Transcript
16:45	☐ Q&A Google panel	Not covered
17:25	☐ Daily Wrap-up Google	Not covered

This edition adds a second attendee's recording of most of the day to my recordings and slide photos. It filled in the welcome Q&A, the end of the robots meta talk, the part on main content, Lightning sessions E and F and the end of the video talk. Still not covered: Lightning session J, "What would you do?", the Q&A and the wrap-up; of Poster session H, only the topics.



A map of Day 2. The processing steps as listed on Google's pipeline slide, with the section of this guide that covers each one.



The three-day arc. One stage of Google Search per day, as shown on Day 1's "How Search works" slide.

02

SECTION 02 • HOW IS HTML INTERPRETED • CHERRY PROMMAWIN

From HTML to extracted elements

Processing starts with the HTML the crawler fetched. Google parses it into a DOM, tells header, navigation and main content apart, and extracts the elements site owners add for indexing.

What the crawler hands over

```
# Illustrative fetch record,
# fields as on the slide
fetch result      ...
connect time      ... ms
time to 1st byte  ... ms
robots policies
  opt-out-google-extended
HTTP response
  HTTP/1.1 200 OK
  <!doctype html> ...
```

SLIDE The crawler passes on more than the HTML: the fetch result, connect time and time to first byte in milliseconds, the robots policies that apply, and the raw HTTP response. (*not in docs*)

SLIDE The robots policies travel with the page into processing; the example showed the Google-Extended opt-out. **DOCS** Google-Extended is a control token, not a crawler (see Day 1).

SLIDE Processing then runs six steps: HTML parsing, rendering, deduplication, feature extraction, signal extraction and index selection (the map in Day 2 at a glance).

What Google extracts from the HTML

ELEMENT	WHAT GOOGLE DOES WITH IT	SEE
rel=canonical	SLIDE In the head; states which page represents another. Used in deduplication and canonical selection.	Duplicates
hreflang	SLIDE Tells Google a page has an equivalent in another language. STAGE Knowing those versions is useful in ranking.	Language
Links	SLIDE "We like links." STAGE Used to discover pages, understand site structure and rank. SLIDE New URLs go back to the crawl queue.	This section
Robots meta	STAGE Probably one of the most important extracted elements, or at least one this audience cares about.	Controls

ANALYSIS

Put **rel=canonical** and **hreflang** in the head of the HTML your server sends, not only in JavaScript-rendered HTML. Google reads them while parsing, without waiting for rendering.

Links Google can and cannot extract

```
<!-- Google can extract -->
<a href="https://example.com/products">
<a href="/products">

<!-- Google cannot extract (slide) -->
<a onclick="goTo('products')">
<span href="/products">
<a routerLink="/products">
<a href="javascript:goTo('products')">
```

SLIDE An `a` element with an `href`, absolute or relative, is the normal way to link. An `onclick` handler never runs because Googlebot does not click, and `href` on a `span` is not a standard link.

DOCS Google's link guide calls these patterns not recommended and notes that Google may still try to parse them.

ANALYSIS The slide is stricter than the docs. Either way, none of them is a link you can rely on for discovery.

STAGE Google sometimes extracts URLs typed as plain text (not in docs; partly unclear in the recording). **ANALYSIS** Don't count on it.

ANALYSIS

Audit every template and component that outputs links: navigation, pagination, filters, product tiles, language selectors. Each needs a real `a` element with an `href`; anything else leaves the target pages without a link Google can rely on.

Four questions from the welcome Q&A

Day 1's leftover questions, answered on slides, then on stage, by Gary Illyes and Cherry Prommawin.

ROBOTS.TXT AND NOINDEX

Why isn't a disallow treated as noindex?

SLIDE "Simply put, it's because of sites that are extremely important and like to disallow their most important pages, either accidentally or out of ignorance." (*not in docs*) Disallowed content isn't indexed, but the URLs can still appear. **STAGE** A national tax authority blocking its key PDFs still gets their URLs shown (*the example not in docs*). **DOCS** A crawler that can't fetch a page never sees its noindex. **ANALYSIS** robots.txt for crawling, noindex for indexing.

HTML SITEMAPS

Do they still help large sites?

SLIDE Good ones are hard to make for large sites: "Better rely on hubs like category pages that link out to your important pages." **STAGE** "if you want to make one, knock yourself out". **DOCS** Google's ecommerce guide: menus to categories to products, plus a sitemap or Merchant Center feed. **ANALYSIS** Google's line has moved since a 2005 post that encouraged HTML sitemaps.

OUT OF STOCK

Keep the page or redirect it?

SLIDE "It depends on the importance of the PDPs to the users." Some users wait months, or pre-order. **STAGE** A rare watch battery is worth the wait; a cheese is simply swapped (*not in docs*). **DOCS** A 302 keeps the source URL in results. **ANALYSIS** Keep the page live with a 200 and current availability in Product markup; redirect only if users would rather switch than wait.

SITEMAPS

List them in robots.txt?

SLIDE "If it's included in the robots.txt file, any crawler can pick your sitemaps up". A sitemap not listed there has to be submitted instead, for example in Search Console. **STAGE** "you can include it in robots.txt. No problem whatsoever." The trade-off: anyone can then see the sitemap.

03

SECTION 03 • JOHN MUELLER, SEARCH RELATIONS TEAM LEAD

Controlling indexing

robots.txt controls crawling, not indexing. Robots meta tags control whether a page is indexed and how it is shown, and some of them now also decide what AI Overviews and AI Mode may use.



"You can use robots meta tags to be more visible in Search results than without robots meta tags."

TRUE OR FALSE? JOHN MUELLER'S OPENING QUIZ. HIS ANSWER: TRUE

SLIDE The tag sits in the head and gives page-specific instructions: `<meta name="robots" content="rule1,rule2">`, or with the name `googlebot` instead of `robots`. **DOCS** Google also respects robots meta tags placed in the body.

Every rule John Mueller covered

RULE	WHAT IT DOES	THE CATCH
<code>all</code>	SLIDE "This is the default value - it does nothing".	It is not an instruction to index the page.
<code>noindex</code>	SLIDE Keeps the page out of results. Typical uses on the slide: internal admin pages, temporary landing pages, thin content.	DOCS Never seen on a URL that robots.txt disallows.
<code>nofollow</code>	STAGE Passes no signals to any link on the page; Mueller called it weird and very broad, and doubted it would exist if the tags were reinvented today.	Use <code>rel=nofollow</code> on the specific links instead.
<code>none</code>	STAGE <code>noindex</code> plus <code>nofollow</code> .	Works only if Google can see the tag.
<code>nosnippet</code>	SLIDE No text snippet or video preview; the title still shows. Also no direct input for AI Overviews and AI Mode.	DOCS A static image thumbnail may still appear.
<code>data-nosnippet</code>	STAGE An HTML attribute that keeps one piece of text, such as a phone number, out of the snippet; it works on a handful of elements.	An unclosed <code>div</code> can hide the rest of the page; valid HTML matters here. DOCS Don't toggle it with JavaScript.
<code>max-snippet:[n]</code>	SLIDE Limits the snippet to <i>n</i> characters; <code>0</code> equals <code>nosnippet</code> . DOCS Also limits use in AI Overviews and AI Mode.	STAGE Too short a value may mean no snippet at all (<i>not in docs</i>).
<code>max-snippet:-1</code>	SLIDE No limit, "can be more than without" a rule (<i>not in docs</i>).	DOCS Google picks the length it finds most effective; no promise of longer snippets.
<code>max-image-preview</code>	SLIDE <code>none</code> , <code>standard</code> or <code>large</code> image previews.	STAGE <code>large</code> matters mainly in Discover.

RULE	WHAT IT DOES	THE CATCH
max-video-preview	SLIDE Seconds of video preview; 0 a static image, -1 no limit.	STAGE He doesn't expect -1 to give an hour-long preview.
notranslate	STAGE No translation offered in results. In a meta tag named <code>google</code> it also turns off Chrome's translation.	DOCS Opts the page out of all translation features in Search.
noimageindex	STAGE None of the page's images are indexed.	Mueller saw no reason to use it.
unavailable_after	STAGE The page drops out of results after a set date and time; most sites don't use it.	DOCS Crawling of the URL slows sharply after the date.
Search generative AI	SLIDE Not a tag but a Search Console setting: exclude keeps links and content out of AI Overviews and AI Mode.	Include is the default; set per property.

SNIPPETS ARE WHAT AI ANSWERS ARE BUILT FROM

STAGE "we use the snippet as a way of building out the AI Overviews and the AI Mode answers". If a page forbids a snippet, Google cannot use it there.

DOCS To appear as a supporting link in AI Overviews or AI Mode, a page must be indexed and eligible to be shown with a snippet.

ANALYSIS Before adding `nosnippet` or a low `max-snippet`, weigh the cost in AI features as well as in classic results.

○ The recommendation for maximum visibility

SLIDE "To be as visible as possible in Google, use:
`max-image-preview:large, max-snippet:-1`"

ANALYSIS Set both on every indexable template unless licensing forbids it, and check that no CMS, plug-in or CDN adds lower limits.

```
<meta name="robots"
  content="max-snippet:-1,
  max-image-preview:large">
```

○ Robots rules, JavaScript and AI crawlers

STAGE Robots meta rules only work if robots.txt lets Google fetch the page. A blocked page's URL can still be indexed; its content can't.

SLIDE Rules can be added with JavaScript, but it takes longer, and removing a `noindex` that way doesn't work. **STAGE** When Google sees a `noindex` in the HTML, "we won't even process the JavaScript". Put robots meta tags in the HTML exactly as intended; use JavaScript only where nothing else is possible, as in a JavaScript web app. **DOCS** Google says it may skip rendering on a `noindex` page. **ANALYSIS** Never serve a `noindex` that a script is meant to lift.

STAGE "AI crawlers don't really know what to do with a `nofollow` link, because they're looking at the content" (*not in docs*). **ANALYSIS** `nofollow` is no access control: use robots.txt per crawler, Google-Extended for Gemini, the Search generative AI setting for AI features.

STAGE robots.txt has an RFC (9309, co-authored by Gary Illies); robots meta tags have none yet, and Mueller invited the room to help write one.


04

SECTION 04 • ERIN SPARLING • LIGHTNING SESSION D

Rendering and JavaScript

Google renders nearly all of the web with a real browser engine, in a detached system that is expensive to run. If content is not in the final DOM, Google can't see it.

SLIDE "Rendering Pages to Index What Users See." **STAGE** In Erin Sparling's words, "our mission is very simple these days: execute JavaScript"; the implementation is complex and expensive (*not in docs*).



The render loop. Redrawn from Google's pipeline slides and the JavaScript SEO basics diagram. Queue time: a community slide said seconds to a couple of minutes; Google's guide says a few seconds, possibly longer, with no upper limit.

DOCS Every page with a 200 status goes to the render queue unless a robots rule says not to index it. **STAGE** Content indexable without JavaScript "can pass through" (*not in docs*). **ANALYSIS** It doesn't wait for the slower pass; the page is still rendered.

STAGE Gemini training reuses Search's rendering: "if it works for Search, it works for Gemini for training", if the site allows training on its content. A page a Gemini user asks about is read live, which takes time, so render fast or on the server (*not in docs*). AI Overviews and AI Mode typically don't read pages live to ground their answers.

Google's four common JavaScript problems

PROBLEM	WHAT HAPPENS	FIX
Content missing from the rendered HTML	SLIDE "The most common issue." (<i>not in docs</i>) Causes: JavaScript Googlebot can't fetch, content waiting for events such as clicks, and disabled browser APIs.	STAGE Server-side or hybrid rendering, fallbacks, no placeholders left in the DOM, differential serving and polyfills. Permission APIs such as location aren't supported.
# in URLs	SLIDE "Fragment identifiers (#) are often ignored by crawlers." A fragment exists only in the browser.	SLIDE "Use the History API for clean URLs in SPAs." STAGE Relatively straightforward.

PROBLEM	WHAT HAPPENS	FIX
Soft 404s	SLIDE The app shows a "Not Found" message but returns HTTP 200.	STAGE A real 404. DOCS Or a JavaScript redirect to a URL that returns 404, or noindex added by JavaScript.
Blocked resources	SLIDE robots.txt disallows critical .js files or API endpoints.	DOCS Block only resources a page doesn't need to be understood.

○ No clicks, no scrolling, a very tall viewport

STAGE Google doesn't scroll; "around 10,000 pixels is what the viewport gets rendered as" (*not in docs*), so content that loads on viewport entry does load. **DOCS** John Mueller called it "viewport expansion" in 2023: rendering a page like a very long phone.

```
// Never runs: Googlebot doesn't scroll
addEventListener('scroll', loadMore);
// Doesn't exist until someone clicks
tab.onclick = () => fetch('/api/specs');

// Fires on viewport entry; it is tall
new IntersectionObserver(load).observe(el);
// In the DOM from the start, hidden
<div class="tab-panel" hidden>...</div>
```

SLIDE "Not visible versus hidden. Present versus absent in the DOM." Hidden content is indexed; absent content is not.

ANALYSIS Treat 10,000 pixels as a margin, not a target. **DOCS** Infinite scroll needs paginated chunks with their own URLs, linked in sequence.

STAGE WebGL effects aren't applied. **ANALYSIS** Feature-detect them, with a fallback.

○ Lightning session D: three community talks

Rendering blind spots · Sören Bendig, Audisto

STAGE "everything you want cited, include it in the raw HTML, server-side rendered". **SLIDE** His five blind spots:

Content changes	STAGE Raw and rendered versions differ: offer counts on a travel site; a bigger newsletter discount in the rendered version than in the HTML the server sent (the two recordings differ on the figures).
Invisible links	STAGE A sub-navigation built only by JavaScript vanishes for bots when rendering fails.
Rendering timeouts	STAGE Slow APIs leave unresolved placeholders in titles and snippets, on different pages each crawl (<i>not in docs</i>).
Content Security Policy	STAGE On a large German banking group's page, a video was blocked by a CSP violation.
Blocked resources	STAGE Never block what rendering needs. Per-session tools that also run for bots can cost a high-traffic site thousands of euros a month (<i>not in docs</i>).

STAGE A page can also render completely blank while its raw HTML looks fine, and one site rendered HTML tags into its meta description, which search engines would ignore (*not in docs*). Check rendering at scale with a modern crawler, not only in DevTools. **DOCS** The renderer may skip requests such as analytics, so client-side analytics won't show its activity fully. **ANALYSIS** Diff raw and rendered HTML per template.

Debugging rendering for Search · Rebecca Yu, JAKALA

STAGE Look for the content in the rendered HTML in URL Inspection: "I always treat this page as the final source of truth". If it's missing, filter DevTools' Network tab by Fetch/XHR, reload, and check that each request runs and isn't blocked or disallowed.

STAGE Four overlooked mistakes: non-crawlable links, hash links included; error pages returning 200 in single-page apps; content behind interaction; blocked scripts or APIs. **SLIDE** A language selector built as a button: "the entire alternate-language cluster is orphaned."

SLIDE "If the renderer can't fetch it, the renderer can't run it." robots.txt applies per host. **ANALYSIS** Check API and CDN hosts too.

```
# Carve out only what
# rendering needs (slide)
User-agent: *
Disallow: /api/
Allow: /api/products/
# Longest match wins
```

Web Fragments · Natalia Venditto, Adobe

STAGE Embedding AI-generated apps can break the host page in three ways: collisions in the global JavaScript scope, CSS bleed, and fate sharing, where the app's crash takes the host down (*not in docs*).

APPROACH SLIDE	JS ISOLATED	STYLES ISOLATED	PART OF THE PAGE	THE CATCH
iframe	✓	✓	✗	"Walled off from DOM, navigation, layout"
Shadow DOM	✗	✓	✓	"Still shares JS globals"
Rewrite everything	By hand	By hand	✓	"Can't rewrite code you didn't write"
Web Fragments	✓	✓	✓	"Needs patches today"

STAGE A hidden iframe runs the app's JavaScript while its DOM sits in a shadow root: "we monkey patch the browser to containerize it". The aim is a standard; ShadowRealm, at stage 2.7, would do much of it (*not in docs*). She recommended server-rendering the fragments: an app reframed that way is fully indexable (*not in docs*). **DOCS** Google flattens shadow DOM when it renders. **ANALYSIS** Check such pages in URL Inspection.

○ Rendering debugging checklist

- ☐ One URL per template in URL Inspection: main content, prices and links are in the rendered HTML
- ☐ Raw and rendered HTML diffed per template; any difference in prices, offers or links is a bug
- ☐ Recurring crawls flag unresolved placeholders (`{{brand}}` , `undefined`); each URL is rendered twice
- ☐ Missing content: robots.txt on every host and the CSP checked via the live test's resources and console
- ☐ Single-page apps: no `href="#..."` routes, and the server returns 404 for unknown paths
- ☐ A valid live test proves access, not indexing; the Page indexing report confirms it

05

SECTION 05 • GARY ILLYES

Understanding what's on a page

Google weighs words by where they sit, stores pages as tokens rather than text, and detects soft 404s with a model that reads layout as well as language.

The main content carries the weight

SLIDE On Google's example blog page, the header and navigation were marked as not so important and the main content as very important: the post title and opening sentence counted; the tagline and the categories sidebar were supplementary.

STAGE Gary Illyes's definition: "Main content is any part of the page that directly helps the page achieve its purpose". It is not only text: images, videos, a tool, user-generated content, the comments below a post, tab content, every heading and the visible title can all be part of it. What sits in the navigation or header tells Google the site doesn't particularly care about it. And "It's the main content that we consider for ranking."

DOCS The Search Quality Rater Guidelines give the same definition; tabs, reviews and comments count as main or supplementary content depending on the page's purpose.



1 Header and navigation

Not so important

2 Main content

Very important: the post title and the opening sentence count

3 Category sidebar

Supplementary, less important text

4 Footer

Unlikely to contribute much to ranking

Where words count. Redrawn from Google's example blog page; the footer weighting was said on stage and is not in the docs.

STAGE Words get different weights depending on the part of the page they sit in: "if you put something in a footer, it's more likely that it's not going to contribute much to ranking" (*not in docs*). The simplest way to make a word count is to move it into the main content, because "where you position text on a page will already contribute quite a bit to ranking that page". And not everything can be main content: then nothing would stand out.

DOCS Google's canonicalization guide calls it the primary content, or centerpiece, and clusters pages whose centerpiece matches. **STAGE** Illyes pointed to the Search Quality Rater Guidelines as the detailed reference. **ANALYSIS** Put the terms a page should rank for in its title, headings and opening paragraphs, not in footer or sidebar blocks.

Tokens, not pages

STAGE The Search index doesn't store sentences or HTML. Text is cut into the smallest units that can still be searched: words, for languages written with spaces; statistically segmented words for Thai and Chinese, with exactly the same segmentation applied to queries (*not in docs*).

STAGE Each token carries metadata for ranking: whether it appeared in the header or the main content ("centerpiece"), in bold or in a heading, plus spam markers such as white text on a white background (*not in docs*). **ANALYSIS** Leftover hidden keyword blocks are a liability, not neutral clutter.

SEARCH TOKENIZER · WORDS STAY WHOLE

tl;dr robots.txt can block bots

AI-MODEL TOKENIZER · SUB-WORD PIECES MAPPED TO IDS

tl ; dr robots . txt can block bots
ID 740

Same text, different units

Search keeps robots.txt and tl;dr whole. The AI tokenizer splits them, punctuation included, and works with numeric IDs, not words.

Two tokenizers. Illustrative sentence; the splits follow Gary Illyes's two slides, where "can" mapped to ID 740 each time it appeared.

STAGE Tokenization for AI models such as Gemini, in training and in inference, differs from tokenization for Search, "or mostly" (*not in docs*). **ANALYSIS** Day 1's slide said Gemini shares tokenization with Search. Read it as a shared step with different outputs.

THE CHUNKING MYTH

STAGE The advice to chunk content for AI is misunderstood. Chunking is real, but it matters at the level of a model's context window, and Gemini's holds millions of tokens (Illyes also put it at 900,000 to a million). Gemini doesn't need 100- or 200-word chunks; thought of in millions of tokens, chunking has perhaps lost its meaning anyway (*that last point not in docs*).


ANALYSIS Don't rewrite pages into short, self-contained chunks. Write for readers: clear headings, complete explanations. Google's long-context docs say 1 million tokens or more: plan with one million as the floor.

How Google spots a soft 404

STAGE A soft 404 should return an error but returns 200, and because error pages are worded in endless ways, keyword matching can't find them (*not in docs*). **SLIDE** Four common causes: pages that look like errors but aren't, thin or empty content, server or CMS misconfiguration, and JavaScript content that fails to load. Plus a bonus: "Mistakes in Google's systems (that you should notify us about)".

STAGE Detection uses a model, "something like BERT, that is specifically trained to understand page structure". Position decides: an error in the navigation need not make a soft 404, but an error alone in the main content does (*not in docs*).

Error: menu failed to load



NOT A SOFT 404

The error sits in the navigation; the main content is real.

Error establishing a database connection

SOFT 404

The main content is only an error, though header and menu look normal.

Position decides. Illustration of the example given on stage.

ANALYSIS Return real status codes for error states, also in single-page apps: 404 or 410 for missing content, 503 for outages such as a failed database connection.

06


SECTION 06 • HANDLING WEB DUPLICATION • JOHN MUELLER • LIGHTNING SESSION E

Duplicates and canonicals

Google clusters equivalent URLs, indexes one representative and forwards the others' signals to it. Which URL represents the cluster is Google's call; your signals only steer it.

STAGE John Mueller gave two reasons to deduplicate: users don't want the same page ten times; and room: the index can't hold everything, storage has many competing uses and its price has risen sharply (*that last point not in docs*). A site migration is simply deduplication across sites: the owner says the old and the new domain are the same, and that Google should pick the new one.

CLUSTERED FROM



One cluster, one representative. Mechanism from Google's duplication slides; the URLs are illustrative.

How clusters form

INPUT	WHAT GOOGLE SAID
Redirects	STAGE Trusted very much: a clear sign there is one version. And "the permanent redirect really only matters for canonicalization, not for clustering." DOCS With a 302, Search keeps showing the source URL. ANALYSIS Use 301 or 308 for moves you want reflected.
Content	STAGE Exact matches (www and non-www), near matches, structurally similar content and soft 404s (<i>the four-way split is not in docs</i>). SLIDE Translated menu and footer around an untranslated main text: "When main content is the same, pages may be clustered."
rel=canonical	STAGE Probably the most common signal, and very often wrong, for example a placeholder value instead of a real URL (partly unclear in the recording). So Google makes its own judgement.
"Magic bucket"	STAGE Everything else.

When a URL pattern decides

SLIDE If `/buy/fax`, `/buy/typewriter` and `/office-equipment` show the same content, Google's systems may assume every `/buy/` URL does: "Do we even need to crawl `/buy/seo-service`?" A car dealer's near-identical city pages can be folded into one city the same way (*not in docs*). The advice: avoid many similar-looking URLs with the same content, and return errors for URLs that no longer exist. **DOCS** To split a wrong cluster, make the pages clearly different; Google may keep them clustered for up to two weeks after the fix.

Bot walls make duplicates

STAGE A CDN challenge page shown to Googlebot on many URLs or domains gets them clustered as duplicates; Google struggles to recognise such pages as errors and can't solve the puzzles. AI agents browsing for users hit the same walls and may buy elsewhere (*not in docs*). **DOCS** Google's CDN post: such pages may be dropped as duplicates and recover slowly; send crawlers a 503.

Choosing the representative

STAGE Many criteria, weighted by machine learning: "we use some kind of machine learning to understand how strong these criteria should be. And this changes from time to time." (*not in docs*)

HIJACKING

SLIDE Several domains claiming the same content: your own staging copy, or third parties, malicious or not. Report cases Google gets wrong.

USER EXPERIENCE

SLIDE The page must load: a broken certificate, failing JavaScript or a page that can't load counts against a URL. The slide also lists meta refresh and security.

SITE-OWNER SIGNALS

SLIDE Redirects, only the preferred URL in sitemaps, and rel=canonical, which "also helps a bit". **DOCS** Redirects and rel=canonical are strong signals; sitemaps weak.

STAGE When every signal points the same way, Google follows the owner; but "if there are multiple things in different directions, we don't know what to do." **ANALYSIS** Before a migration or template change, align redirects, rel=canonical, internal links, sitemaps, hreflang and working HTTPS.

Same language, different countries

SLIDE German pages for Germany, Austria and Switzerland are tricky to deduplicate. "We try to use hreflang alternates." Avoid "clever" geo-redirecting. **DOCS** Pages in hreflang clusters are preferred as canonicals. **ANALYSIS** Use de-DE, de-AT, de-CH and make the pages differ in more than boilerplate.

Old domains after a migration

SLIDE Alternate names are why `site:` queries still show an old domain, which owners misread as a failed move. Someone searching an old brand may see the old URL. **STAGE** The redirect takes them on. **DOCS** "This is normal and as users get used to the new domain name, the alternate names will fade away without you doing anything." **ANALYSIS** Judge the move in Search Console; keep redirects long term.


STAGE Pointing a paginated series' canonicals to page 1 can sometimes make sense, for example to strengthen a category page (*not in docs*). **DOCS** Google's 2013 post calls it incorrect: later pages' content would not be indexed.

ANALYSIS Keep self-referencing canonicals unless you deliberately want later pages folded into page 1.

SLIDE The closing list: use redirects for migrations, use HTTP result codes, "Don't block agents.", check your rel=canonical links, use hreflang to help Google localise, report weird canonicals in the forums, make secure pages that work, and keep canonical signals clear.

Canonical graphs • Tobias Schwarz, Audisto

STAGE In this community talk, checking canonicals pair by pair misses the picture. A canonical group is every URL connected by canonical links; its leader is the page they all lead to. Audit whole groups, with redirects and the indexability of each member: broken canonicals let the wrong pages show up in results.




Canonical groups and their leaders. Illustration of the patterns shown in the community talk; the real-world chain ran from a feed URL through a 301 on another subdomain.

PATTERN	COMMUNITY SPEAKER STAGE	GOOGLE'S DOCS DOCS
Several canonicals on one page	Invalid; the engine falls back on its own heuristic, or ignores the tags completely.	Google will likely ignore all of them.
Chain	Improper use; runs through canonicals or through server- or client-side redirects.	Google can follow chains but recommends one canonical target. ANALYSIS A gap in severity, not in advice.
Loop	No leader, so the group's canonical information can't be used (<i>not in docs</i>).	—
Leader users can't reach	"you're telling the search engine to index something that a human can't reach". Make the linked page the leader.	—
Canonicals across languages	Misuse: only one language version would rank. Connect versions with hreflang.	On hreflang pages, the canonical should be in the same language.
Weak leader	Valid but likely suboptimal; the strongest page should lead (<i>not in docs</i>).	Links to duplicates are consolidated into the canonical. ANALYSIS The real risk is Google choosing another page.

ANALYSIS To audit at scale, resolve every crawled URL to where its canonicals and redirects finally lead, group URLs by that leader, and flag groups with more than one hop, a loop, several canonicals on a page, or a leader that is noindexed, blocked, non-200 or redirecting.

○ Merging two competing sites • Martyna Aġanoġlu, Loando

STAGE Two of Poland's first loan-comparison sites did exactly the same thing and competed for the same users, keywords and results. After an acquisition by a group active in eleven markets, they became one brand. Three goals came before any technical work: lower costs, one strong brand, and a simple structure instead of hundreds of similar URLs competing with each other.



Keep, merge or remove. The decision every URL got in the community case study, with the speaker's figures.

STAGE Both domains were checked for past penalties and compared on traffic, conversion and revenue before one was chosen. Removed URLs got an error status; none went to the homepage, even those still bringing traffic. Then: new sitemaps, a change of address in Search Console, internal links pointing at the new pages, merged analytics, the financial experts behind the site introduced, and Search Console checked daily against the old baseline for months. **DOCS** Google's site move guide lists the same three technical steps.

STAGE Traffic grew from the first day, with no dip; revenue grew 100% month over month (the speaker's figure, period not stated). **DOCS** Don't redirect many old URLs to one irrelevant page such as the homepage; return 404 or 410 for content that isn't moved. **ANALYSIS** Which of the two codes matters less than not redirecting to an unrelated page. At about 2,000 URLs crawl budget is hardly the gain; one strong URL per intent, with its signals consolidated, is.

○ A migration after an acquisition • David Carrasco Pamies, Magnify



"I think there are no complex migrations. There are complex people."

DAVID CARRASCO PAMIES, ON A MIGRATION STILL UNDER WAY

ABSORB	KEEP APART	MERGE	PLUG IN
STAGE One site is integrated into the other	Both sites stay	Two or three sites become a new one	A full site becomes a section of the other

STAGE Run the migration around one document, the redirect map, which gives every old URL an approved destination. Most failures look technical but trace back to a decision nobody made: in his project, more than 7,000 URLs still had none. Prioritise by business value, not only traffic; ask support and sales which pages make buyers buy, and other teams about what a crawler can't reveal, such as legal requirements.

STAGE Before launch, freeze the inventory: CMS export, crawl, sitemap, search data and logs. After it, 200s across the new domain prove only that URLs work. Test every URL's route, answer and lead, automated, with people checking the critical journeys, and give moved content its own baseline.

STAGE An unnamed study's median recovery is 304 days, but broken routes need action at once. **ANALYSIS** Google's guide says weeks until new URLs show; traffic takes longer: report both.

```
# One row of a redirect map,
# fields as in the talk (illustrative)
old URL    /pricing-enterprise
new URL    /products/x/pricing
survives   the product answer,
           the demo request
approves   the decision's owner
test      route • answer • lead
```

07

SECTION 07 • FINDING THE GOLD NUGGETS • LIGHTNING SESSION F

Feature extraction: structured data, images, video

After deduplication, Google pulls the "gold nuggets" out of each page: structured data, images and videos. Markup is still the precise, cheap route, and Google says it will keep preferring it.

STAGE A dedicated system, whose internal name Gary Illyes wouldn't share, does the expensive extraction, and it runs after deduplication so the work is spent on a trimmed set of documents (*not in docs*). **DOCS** Google's guidelines: "we recommend placing the same structured data on all page duplicates, not just on the canonical page". **ANALYSIS** Markup that exists only on a duplicate that loses canonical selection may never be extracted.

Why structured data still matters

STAGE Ryan Levering, a Google software engineer who works on data ingestion, described two camps inside Google: markup is useless because models can read pages, or markup is the future of machines talking to each other. The truth, he said, is in the middle. **SLIDE** Four reasons it still matters:

PRECISION

SLIDE Higher accuracy than large-scale LLM extraction for complex schema such as sale pricing. **STAGE** In his tests LLMs invented properties and got nested pricing wrong (*not in docs*).

EXTRA CONTENT

SLIDE Metadata the page text lacks, such as full ISO dates (Google may misread a visible date's time zone) or stable IDs. **DOCS** Still only for content users can see.

EFFICIENCY

STAGE Even Google can't run complex models on every page; parsing markup is nearly free: "So we're always going to prefer that particular approach." (*not in docs*)

FOCUS

SLIDE Markup points at the pertinent data, so Google doesn't pick up, say, a related product's price (*not in docs*).

STAGE The same processed markup feeds classic results and AI Overviews and AI Mode. Raw schema.org is generally not put straight into a model's context: it is sorted, checked and indexed first, partly because raw markup would be an abuse vector (*not in docs*). **DOCS** Google's guide to generative AI features lists overfocusing on structured data among things not needed; it still helps rich results. **ANALYSIS** Skip markup written for AI; implement the types Google documents.

What Google recommended

- STAGE** JSON-LD, microdata and RDFa enter the same pipelines; JSON-LD is recommended because people make fewer mistakes with one block. Microdata saves page weight (*not in docs*).
- STAGE** "we will never penalize you for having more structured data on your pages", but irrelevant markup can be treated as abuse. **DOCS** A manual action removes rich results, not rankings.
- STAGE** A common fault: two plug-ins emit the same type, and an event page reads like a list of events (*not in docs*). **ANALYSIS** Test one page per template.

2.7×

more search traffic for Rakuten Recipes after adding recipe markup, plus 1.5× session duration

Case study, May 2018; "old", said the speaker

- **STAGE** Use unique identifiers (an organisation's homepage URL) and only relevant types. **SLIDE** Find them in the Search gallery, and test in the Rich Results Test first.

Coming and changing

SHACL validation rules

SLIDE Coming soon: downloadable rules on each feature guide, to check markup in your own build before publishing (*not in docs*). **STAGE** Search Console stays the reference.

Merchant markup

STAGE Sale prices can carry validity dates; loyalty programs and shipping policies were added in 2025. **DOCS** The July 2026 docs log calls the sale-dates change a clarification.

Adoption decides

STAGE schema.org now publishes usage statistics per type, because Google often waits for adoption before building a feature (*not in docs*). Little-used features were removed in 2025.

Images and video

ON THE PAGE	WHAT GOOGLE SAID
<code></code>	STAGE What Google's HTML parser looks for. Extracted images go to the image indexing service and can appear in Discover, Images, web results and AI features.
<code><picture></code>	STAGE Supported only because it must contain an <code>img</code> , which is what gets extracted.
CSS background	STAGE "we don't support CSS background extraction"; Illyes even suggested it to keep an image out (<i>not in docs</i>). DOCS The documented ways are a robots.txt rule for Googlebot-Image or a noindex X-Robots-Tag.
src and alt	STAGE Google uses three <code>img</code> attributes (<i>not in docs</i>): <code>src</code> matters most, then <code>alt</code> , which should describe the image for someone who can't see it, and lets it rank for those concepts. DOCS The image guide calls alt text the most important attribute for image metadata.
Surrounding text	STAGE Critical for understanding and ranking the image. ANALYSIS Give key images a caption, not just alt text.
Image sitemap	STAGE The other reliable way to tell Google about images.
Formats	STAGE Nearly all popular formats work; AVIF currently has ingestion hiccups (<i>not in docs</i>). DOCS AVIF is listed as supported. ANALYSIS Keep WebP or JPEG in the <code>img src</code> .
Video	SLIDE A dedicated watch page per video, listed in a sitemap. STAGE Without a watch page, or below the fold, a video won't be indexed (<i>the fold part not in docs</i>); the HTML needs a video container, and JSON-LD on top. DOCS Only watch-page videos are eligible, with the player on the page at load, not behind a click.

SLIDE The seven video factors also include high-quality content, compelling titles and descriptions, relevant thumbnails and fast pages, and a warning: "just because you spent money on a genAI subscription, that does not mean you are going to do well with AI generated videos"

The rest of the video talk

- **STAGE** A wrong thumbnail makes very many viewers drop out at the start. "Fast-loading pages" really means a fast-loading video: a CDN that beats your own server is a win (*not in docs*).

Robots rules for images and video

- **STAGE** A media indexer attaches images and videos to their host page's URL (*not in docs*). Googlebot-Image and Googlebot-Video fetch them; a robots.txt disallow on the file keeps it out.

- **STAGE** The popular formats are supported; Google recommends the MP4 container and standard codecs (*not in docs*).
- **STAGE** Video sitemaps aren't critical, but Google ingests them far more often than it can process HTML, and they say which URLs to check (*partly not in docs*). Text around the video, and people talking about it, help it surface.
- **STAGE** Consider YouTube or Vimeo: they solved video search years ago.
- **ANALYSIS** Put a watch page's main video in the first viewport and load the player without a click-to-play placeholder.
- **STAGE** Unlike a blocked web page, a blocked video is never shown by its URL alone (*not in docs*).
- **STAGE** `noimageindex` also stops videos, whose thumbnails are images (*not in docs*). **ANALYSIS** Never set it on watch pages.

AI-generated images

- **STAGE** Up to the site owner. Google shows AI images "if they are looking for them specifically" (*not in docs*). Diffusion models garble text: Google's own 2025 timeline image jumped from 1994 to 1995 and back. Check them for hallucinations and regenerate.
- **DOCS** Google Images reads IPTC `trainedAlgorithmicMedia` and can show C2PA details in "About this image".

○ Lightning session F: two community talks on media

A COMMUNITY SPEAKER

LLM alt text that missed the point

STAGE Since version 20 (May 2024) Screaming Frog can call ChatGPT through custom JavaScript, so SEOs rewrote alt text at scale. For a page selling an anxiety medicine for cats, the model saw only a cat and a veterinarian and dropped the page's intent; image performance in Search Console fell after the rollout. Test in List mode on chosen URLs first. A safer script adds the page title to the image description; it still covers only the hero image and one market. **DOCS** Google says AI-written alt text should be fact-checked before publishing.

PATRICK DOMANICO • DENTSU

"Prompt journalism": more from what you recorded

STAGE Optimise interviews and podcasts you already have, not more content. Clean the transcript of fillers with editorial judgment, without correcting the speaker's grammar (W3C guidance). Let AI map it into chapters, headings and entities, which can feed metadata and structured data; write the inverted pyramid as a prompt for the article, then cut angles for LinkedIn, Substack, X and YouTube. Social accounts can be verified in Search Console to track that content.

08

SECTION 08 • INTERNATIONALISATION • LIGHTNING SESSION G • POSTERS

Language, country and hreflang

Google reads a page's language from its content and records one language per page. hreflang only maps equivalent versions to audiences, and it fails on small mistakes.

STAGE Working out a site's localised pages and target countries is one of the complex tasks of indexing, Google said, and started with the basics: Google may find language versions on its own, but each needs its own URL; switching languages by cookie or by swapping content in place makes them very hard to handle. **DOCS** Googlebot mostly crawls from US IP addresses and sends no Accept-Language header.

```
<!-- In the head of every version,
      the page's own URL included -->
<link rel="alternate" hreflang="en"
      href="https://example.com/en/">
<link rel="alternate" hreflang="en-gb"
      href="https://example.com/en-gb/">
<link rel="alternate" hreflang="es"
      href="https://example.com/es/">
```

The mistakes Google sees most

- **SLIDE** "If page X links to page Y, page Y must link back to page X." **STAGE** One-way pairs are ignored, or anyone could claim to be your language version.
- **STAGE** No self-reference: every page must list its own URL too.
- **STAGE** Two or three methods at once, conflicting. **DOCS** HTML head, HTTP header and sitemap are equivalent; pick one.
- **DOCS** Alternate URLs must be fully qualified, protocol included.

hreflang code mistakes

STAGE Wrong language and region codes are a mistake Google finds very often, the same ones again and again, especially in Europe, because owners assume the codes are easy.

WRITTEN	MEANT	CORRECT	WHY SLIDE STAGE ANALYSIS
se, se-SE	Swedish (Sweden)	sv , sv-SE	The language part is ISO 639-1; SE is Sweden's country code. se is Northern Sami.
dk	Danish	da	DK is Denmark's country code, not the language.
cz	Czech	cs	CZ is Czechia's country code and domain; Czech is cs .
en-UK	English (United Kingdom)	en-GB	Regions are ISO 3166-1 Alpha-2 codes.
de-SW	German (Switzerland)	de-CH	Switzerland is CH.
de-GE	German (Germany)	de-DE	GE is a valid code, for Georgia, so the error is silent.
en-EU	European Union	One per country	hreflang targets countries only.
en-LA, en-SA	Los Angeles, South Africa	A country code	LA is Laos and SA is Saudi Arabia; a city is never a region.
GB	United Kingdom	en-GB	A region code can't stand alone; the value must name a language.

○ Language comes from the content

STAGE Google doesn't use hreflang to determine a page's language, nor a language code in the URL: it reads the content, and annotates each page with one language (*that part not in docs*). Pages that mix languages are hard to place.

DOCS Pages that translate only the menu and footer around the same main text count as duplicates. **ANALYSIS** Translate templates together with the main text.

○ Country targeting

STAGE A ccTLD is one of the strongest country signals; server location (an unclear word in the recording) is not really reliable nowadays (*not in docs*). **ANALYSIS** A second recording heard "IP address" for the strongest signal; the docs back the ccTLD. **DOCS** Google still lists server location as a possible, not definitive, signal, and marks URL parameters such as `?loc=de` as not recommended. **STAGE** ccTLD, subdomain or subdirectory are all fine; the choice depends on availability, cost, local law and commitment. "Don't try to move your domain now."

DOCS Google's locale signals also include local addresses and phone numbers, local language and currency, links from other local sites and Business Profile signals.

AI ANSWERS HIDE LANGUAGE GAPS

STAGE AI Overviews and AI Mode need no extra work: they sit on top of web results. But in classic results users notice missing or poorly translated content; AI answers synthesise it away, "an invisible gap". And "AI is still language-dependent": a query in another language draws on different data.

ANALYSIS Test AI answers with native-language queries in each market, and compare with the country breakdown of Search Console's generative AI report.

STAGE Machine translation is the owner's call, weighed against quality, local conventions and culture; dates and calendars differ (in Thailand the current year is 2569) (*not in docs*). **DOCS** Automated translation of scraped content is a listed example of scaled content abuse. **ANALYSIS** Treat machine translation as a first draft for a native reviewer.

○ Lightning session G: two community talks

ALIZÉE BAUDEZ • INTERNATIONAL SEO CONSULTANT

The translation was right; the region wasn't checked

STAGE A furniture maker's Dutch site converted at a fraction of its other markets despite clean translations and hreflang; Dutch homes are narrow, with steep stairs. "the translation was right, but the region was never checked". hreflang for de-AT doesn't make a brand trusted in Austria. If you can't name three local sources that see your brand as present: "don't bother expanding to another country; just fix the ones you already have". She cited a 2017 Semrush study: 75% of 20,000 multilingual sites had at least one hreflang error.

COMMUNITY TALK • PRESENTED BY A COLLEAGUE

Non-Latin scripts: letters, direction, keyboards

STAGE Arabic and Persian have letters that look identical but are different Unicode characters, such as the Arabic and the Persian ye, so a searcher may type one variant while the site uses the other, and keyword data splits across variants (*not in docs*). Titles, product names and URLs that mix right-to-left and left-to-right text can display in a confusing order. And the same Persian query may be typed in Latin letters. **DOCS** Google's 2023 post gives the same case for Hindi.

ANALYSIS Research each spelling separately and add up the volumes.

STAGE In competitive Persian, Turkish and Arabic searches, bought links and paid editorial content still visibly sway rankings, the presenter observed, as a description and not a recommendation (no data shown) (*not in docs*); Google treats buying links as spam and in October 2023 improved spam coverage by language. **DOCS** That update names Turkish and "other languages", and cloaking, hacked, auto-generated and scraped spam. **ANALYSIS** It names neither Persian nor link spam, so it doesn't show the bought links were dealt with.

STAGE AI answers to Persian queries are written in Persian but cite some English sources, because thinner Persian content sends retrieval to other languages ("cross-lingual retrieval"). Features also launch language by language. **DOCS** Site names reached all languages only in September 2023. **STAGE** Persian offers fewer natural link opportunities; base the strategy on each market's maturity. Multilingual SEO is not just translation.

○ Poster session H

STAGE Posters were added for topics that need a longer, one-to-one discussion than a seven-minute talk. Topics included technical SEO testing in CI/CD pipelines, internal links that move rankings, the URL Inspection tool and API, bringing everyone on board in complex migrations, moving a CMS from WordPress to Contentful, the "global Spanish" problem and going beyond prompts. One poster, by a Googler, was on bringing Google's tools into WordPress sites.



09

SECTION 09 • CALCULATING SIGNALS • INDEX SELECTION

Signals and index selection


Some signals are calculated at indexing because ranking has no time for them. Index selection then applies thresholds to decide what the finite index keeps, and quality decides most of it.

○ Signals calculated at indexing

Two talks by the same Google speaker.

SIGNAL	WHAT WAS SAID STAGE	DOCS AND ANALYSIS
Country and language	"Both of them are extremely, extremely important". Language is weighted in index selection, otherwise "you would have an English index"; country too, so the UK or Switzerland isn't swamped by the US or Germany (<i>not in docs</i>). Both are used again in ranking.	ANALYSIS A small country in a big language should make its country unmistakable.
Freshness	Matters for queries that deserve fresh results, such as a breaking event; partly calculated at indexing, used in ranking. "It doesn't play a big role in index selection." (<i>not in docs</i>)	ANALYSIS Changing a page's date won't get it indexed.
SafeSearch	Calculated at indexing from the whole page and its links, because "In ranking, everything has to happen online, and there's just not enough time to calculate those things." Explicit content doesn't lower indexing chances (<i>not in docs</i>).	DOCS Keep explicit pages on a separate domain or subdomain.
Spam	"If we detect that something is spam, then we're not letting it into the index." SpamBrain is now built on Gemini (<i>not in docs</i>); over 99% of visits from Search are spam-free.	DOCS SpamBrain launched in 2018; it found 5× more spam sites in 2022 than in 2021. ANALYSIS Quote these documented figures.

ANALYSIS Because these signals are fixed at indexing, a fix (a page's language, content that triggers SafeSearch) only counts once Google recrawls and reprocesses the page.



Where pages stop. Two Search Console statuses mapped to the stage that produces them. In the dropped list, "noindex" is the likely reading of one unclear word in the recording.

○ How index selection decides

- **STAGE** "our index is immense. Like, immense. But it is a finite resource." Most URLs wouldn't be useful to users, so thresholds decide what is kept.
- **STAGE** It is a predictive system built on machine learning, and it uses what Google knows about the site: "the index selection system is going to be more forgiving when it sees a new URL from your site" if the site or section serves users well (*not in docs*). **ANALYSIS** Launch new pages where Google already indexes well.
- **STAGE** Where coverage of a country or language is thin, the bar drops: in a language such as Basque, a site that starts publishing can very likely replace spam in the index (one condition was inaudible) (*not in docs*). **ANALYSIS** An opening for content written for that market, not for bulk machine translation.
- **STAGE** Importance counts: news sites' pages usually get indexed (a likely reading) very fast (*not in docs*). In the end "focusing on the quality is the most reliable way to get stuff in the index".
- **STAGE** Non-canonical duplicates are dropped, so "only canonicals" reach results. **DOCS** An alternate can still be shown in some contexts, such as a mobile URL to a mobile user.

○ Reading the Page indexing report

STATUS	ON STAGE STAGE	HELP PAGE DOCS	WHAT TO DO ANALYSIS
Discovered – currently not indexed	A crawl-scheduling state: Google doesn't want to crawl it yet; "The first one is kind of nastier." Getting other pages indexed and showing they're useful helps (<i>partly not in docs</i>).	Found, not crawled, typically because a crawl was expected to overload the site.	Rule out slow responses and server errors, then raise quality; resubmitting changes nothing.
Crawled – currently not indexed	An index selection decision; "most of the time it is actually a quality issue" (<i>not in docs</i>). If quality is even, look at a different template that hides the content.	May be indexed later; no need to resubmit.	Compare with indexed pages of the same type: thin, duplicate, soft-404-like content, template.

STAGE Use the report's not-indexed reasons to test changes. And rather than hunting signals in the docs: "If you are just focusing on actually creating the content that users will like, then that's probably a better use of your life."

10

SECTION 10 • HOW DOES THE INDEX LOOK LIKE?

Inside the index

The index holds tokens with their metadata and signals, not pages. Retrieval runs on posting lists and on embeddings, and AI Overviews and AI Mode use the same index and the same snippets.



Two ways in, one index. Illustrative query and document numbers; the mechanism as described on stage, which Google has not documented in this detail.

STAGE Each page's tokens are stored as they are, with their metadata: "we don't have the full content of the page in our index" (*not in docs*). Nearly every signal calculated for a document, such as quality, country and language, is attached to it.

STAGE Posting lists, a structure at least 60 years old, list for most tokens the URLs that contain them; the lists of a query's important words are intersected into unranked candidates (*not in docs*). **DOCS** Google's explainer simplifies this to "an entry for every word".

STAGE "just like you build a posting list, you can also build a vector space": the distance between embeddings of document parts and of the query decides what comes back. Both routes work from page content, so "this mantra that we started 25 years ago or whatever still stands, whether we like it or not".

THE SAME INDEX, THE SAME SNIPPETS FOR AI ANSWERS

STAGE "the snippet that you see was reconstructed from these tokens" (*not in docs*). AI Overviews and AI Mode use those same structures and snippets: fan-out queries go to the Search index and come back with snippets that feed the answer. **DOCS** nosnippet and max-snippet govern that use.

ANALYSIS If a page never appears as a source, check indexing, snippet rules and the Search generative AI setting first. Name the subject in the words people search with and explain it fully; keyword variations add nothing.

11

SECTION 11 • OMRI WEISMAN, ENGINEERING MANAGER, GOOGLE TRENDS

Google Trends for SEO

Google Trends is a sample of aggregated, anonymised searches, shown as search interest, not volume. Google presented three uses: keyword selection, scheduling and ideation.

STAGE Trends covers web, image, shopping and news search plus YouTube, but not Discover, by world, country or region, "from 2004 and up until three minutes ago". **DOCS** Internal searches made by AI Mode and AI Overviews are filtered out, and privacy noise means a one-off spike on a rare query may not be real.

Interest is not volume, and spikes need context

STAGE Interest in ski has declined since 2004, yet Google's internal numbers show ski searches rising (*not in docs*): "Search interest is not search volume." **ANALYSIS** A falling line is not a shrinking market; take volumes from Search Console or keyword tools.

STAGE An early-2012 spike for "ai" came from a hit song with the Portuguese word in its title, not from artificial intelligence (*not in docs*). **ANALYSIS** Read the queries behind a spike before acting on it.

STAGE Worldwide interest in Gemini passed ChatGPT for the first time in September 2025, in the Nano Banana surge, and its baseline stayed higher (*not in docs*). **DOCS** Nano Banana was announced on 26 August 2025. **ANALYSIS** Judge a launch by the baseline after the spike.

Three uses for SEO

USE	HOW GOOGLE SHOWED IT STAGE	TIP
Keyword selection	Compare terms on the Explore page for the right country, region and time frame; use the higher-interest wording; check top and rising queries.	ANALYSIS Compare per market: the winning wording can differ within one language.
Scheduling	Seasonal patterns predict the next wave. In the US, umbrella and sunscreen searches both peak in June (<i>not in docs</i>).	ANALYSIS Publish a few weeks before the rise in each target country.
Ideation	Trending now shows a real-world event within about ten minutes; journalists check it every morning for gaps.	DOCS "you shouldn't write about something just because it's trending"

New in Google Trends

A new Explore page

STAGE Launched early 2026; the legacy page will be retired (*not in docs*). "Suggest search terms" uses Gemini to pick up to eight terms from a plain-language request.

Comparisons, maps, categories

STAGE Year-, month- and week-over-week chips; regional maps; the category filter is back (*the last two not in docs*). **PRESS** Search Engine Roundtable, August and September 2026.

Trending now, Trends TV

STAGE Trending now covers 125 countries, and trends.google.com/tv suits an office screen. **DOCS** The help page says 100+ countries and regions.

12

SECTION 12 · ACROSS THE DAYS

What changed since Day 1

Day 2 mostly explained the mechanisms behind Day 1's slides. Two points pulled against what Day 1 said or Google documents.

TOPIC	DAY 1	DAY 2
Tokenization Contradicts	SLIDE Gemini shares technologies such as tokenization with Search.	STAGE The two tokenizers differ, "or mostly". ANALYSIS A shared step with different outputs.
Paginated canonicals Contradicts	DOCS Don't use page 1 as the canonical of a series.	STAGE It can sometimes make sense. DOCS The 2013 post still calls it incorrect.
A finite index	SLIDE Indexing is one big but not limitless index.	STAGE Immense but finite; storage is contested and dearer; a predictive ML system decides what is kept.
AI on the same index	SLIDE AI Mode and AI Overviews use Search's crawl and index; fan-out expands the query.	STAGE They are built from the same token-based snippets; fan-out queries hit the same index.
Chunking	SLIDE No need to chop content for AI.	STAGE Chunking matters only at the context window, which for Gemini holds millions of tokens.
The pipeline	SLIDE Fetch replies go to indexing; changes in connect time and TTFB (and 429/5xx) drive hostload.	SLIDE The fetch record carries those timings and the robots policies; processing has six steps.
No clicking	DOCS Crawlers don't click buttons. STAGE "Load more" is risky.	STAGE Click-loaded content isn't supported; no scrolling; a viewport of about 10,000 pixels.
Soft 404s	DOCS A success code on an error-like page: excluded, still crawled. SLIDE BERT reads words in context.	STAGE Found by a BERT-like model reading layout; common in single-page apps. After a migration, 200s prove only that URLs work.
Bot challenges	ANALYSIS Never serve a challenge page with a 200.	STAGE Such pages get clustered as duplicates. DOCS Send crawlers a 503.
Inheritance by path	SLIDE Unknown URLs inherit their parent path's quality and popularity.	STAGE Index selection is more forgiving for strong sections. SLIDE A URL pattern can make crawling a new URL unnecessary.
robots.txt and noindex	ANALYSIS A blocked URL can still be indexed, and its noindex is never seen.	SLIDE Why Google won't treat a block as noindex: important sites block their best pages by mistake. STAGE An important URL may be indexed without its content; robots meta rules need a fetchable page.
Opting out	STAGE robots.txt is how site owners opt out of Google's crawling.	STAGE Google strongly believes owners should be able to opt out, for legal reasons or for crawl budget.
Gemini	SLIDE Gemini uses crawlers for data and grounds on the Search index.	STAGE Its training reuses Search's rendering (where training is allowed); SpamBrain is now built on Gemini.
Repeated	SLIDE 5 trillion+ searches a year. DOCS The generative AI opt-out.	STAGE The same figure, announced in early 2025. SLIDE The opt-out, again.

13

SECTION 13 • TAKE IT BACK TO THE OFFICE

Day 2 action checklist

Seventeen checks that turn Day 2 into work. Most need URL Inspection, a crawler and a look at your templates.

RENDERING AND INDEXING CONTROLS

- ☐ Title, canonical, robots meta, main text and links are in the HTML the server sends
- ☐ Tab content is in the initial DOM, lazy loading fires on viewport entry, long lists have paginated links
- ☐ robots.txt on every host (site, API, CDN) allows the scripts and endpoints rendering needs
- ☐ Pages to keep out are crawlable and carry `noindex` in the server HTML
- ☐ `max-snippet:-1` and `max-image-preview:large` on indexable templates; snippet limits weighed as AI limits

DUPLICATES, CANONICALS AND SITE MOVES

- ☐ Redirects, `rel=canonical`, internal links, sitemaps and `hreflang` all name one preferred URL
- ☐ Canonical groups: one hop, no loops, one canonical per page, an indexable and linked leader
- ☐ Bot challenges reach crawlers as a 503; empty variant combinations return 404
- ☐ Every moved URL has an approved same-intent destination; removed URLs return 404 or 410, never the homepage

CONTENT, MARKUP AND MEDIA

- ☐ Ranking terms sit in the title, headings and opening text; no hidden keyword blocks remain
- ☐ One page per template passes the Rich Results Test without duplicate types; duplicates carry the same markup
- ☐ Images that should rank use `img` with a real `src`; key videos have their own watch page and markup
- ☐ Watch pages load the player in the first viewport without a click, and carry no `noimageindex`

INTERNATIONAL

- ☐ `hreflang`: self-reference, return links, one method, codes validated against ISO 639-1 and 3166-1
- ☐ Templates are translated with the main text; the language selector uses plain `<a href>` links

INDEX SELECTION AND MEASUREMENT

- ☐ Not-indexed URLs are compared with indexed pages of the same type, not resubmitted again and again
- ☐ Trends lines are read as shares; volumes come from Search Console or keyword tools

14

SECTION 14 • BEYOND THE DOCUMENTATION

Said on stage, not in Google's docs

Points from Google's speakers that I could not find in Google's documentation. Not wrong for that, but not commitments either: treat them as informed hints and recheck before you build on them.

POINT	WHO	WHY IT MATTERS
STAGE Footer words weigh less	Gary Illyes	ANALYSIS Ranking terms belong in the main content.
STAGE The index stores tokens and metadata, spam markers included	Gary Illyes	ANALYSIS Hidden text is recorded, not ignored.
STAGE Snippets are rebuilt from token positions	Google	AI Overviews and AI Mode use the same snippets.
STAGE A BERT-like model reads page structure to detect soft 404s	Gary Illyes	An error alone in the main content is enough.
STAGE The rendering viewport is about 10,000 pixels tall	Erin Sparling	Lazy loading on viewport entry works.
STAGE Gemini training reuses Search's rendering (where training is allowed)	Erin Sparling	ANALYSIS No separate rendering work for Gemini.
STAGE Machine learning weighs canonical criteria; SLIDE URL patterns can decide clusters	Google	The weighting changes from time to time.
STAGE Index selection is predictive, and forgiving for strong sections	Google	ANALYSIS Launch new pages where Google already indexes well.
STAGE Language and country balance the index; freshness and explicit content don't sway selection	Google	Thinly covered languages get a lower bar.
STAGE SpamBrain is now built on Gemini	Google	Google says its AI spam detection is highly accurate.
STAGE Google will always prefer markup to model extraction	Google	ANALYSIS Keep markup for prices, offers and dates.
SLIDE SHACL validation rules for structured data are coming	Google	ANALYSIS Plan an automated markup test.
STAGE 'Discovered – currently not indexed': Google doesn't want to crawl it yet	Google	DOCS The help page cites expected server overload.
STAGE A video below the fold is not indexed	Gary Illyes	ANALYSIS The docs only require the player at load; put it in the first viewport anyway.
STAGE <code>noimageindex</code> also stops videos; a blocked video never shows as a bare URL	Gary Illyes	Media controls reach further than their names suggest.

15

SECTION 15 • STAY IN TOUCH

Help complete this guide

This guide covers what I could capture in the room. If you heard something on Day 2 that isn't written here, tell me and I'll add it, so everyone who attended ends up with the full picture.

01 • ADD WHAT I MISSED

Heard something that isn't here?

A second attendee's recording filled most of the first edition's gaps. These are still missing:

- Lightning session J: Indexing and Troubleshooting
- What would you do?: Indexing Edition
- Poster session H: everything beyond the topics
- The Q&A and the daily wrap-up

Each addition is labelled and checked against Google's documentation before it goes in, and credited to you by name if you like.

02 • THE FULL KNOWLEDGE BASE

Every Day 2 point, in the web edition

This guide is a selection. The web edition holds all of Day 2's points, nearly 1,000, with topic pages, session pages, a search and the verification status of every point. It is a folder that opens offline from [index.html](#).

An agent-ready version of the same data lets AI tools answer questions about the event, with every answer traceable to its slide, session or document.

Message me and I'll share both.

LET'S KEEP IN CONTACT

Ibrahim Anjro

TEDX SPEAKER

WHATSAPP +49 1575 1702770

LINKEDIN [linkedin.com/in/ibrahem-angro](https://www.linkedin.com/in/ibrahem-angro)


Message me with additions, corrections, or to get the web edition.



WHATSAPP



LINKEDIN



REFERENCES

Sources

This guide is built from 64 photos of the slides, my recordings of 16 sessions and a second attendee's recording of most of the day, transcribed and checked against the slides where I had them. Every point labelled Docs was checked against these pages on 3 October 2026.

1. Search generative AI control — Google Search Console Help
2. Optimizing your website for generative AI features on Google Search — Google Search Central
3. Introducing Search Generative AI performance reports in Search Console — Search Central blog (3 June 2026)
4. In-depth guide to how Google Search works — Google Search Central
5. Google Search Essentials — Google Search Central
6. A guide to Google Search ranking systems — Google Search Central
7. Creating helpful, reliable, people-first content — Google Search Central
8. List of Google's common crawlers — Google
9. How Google interprets the robots.txt specification — Google
10. Optimize your crawl budget — Google
11. How HTTP status codes affect Google's crawlers — Google
12. Troubleshoot Google Search crawling errors — Google Search Central
13. Qualify your outbound links to Google — Google Search Central
14. Evolving "nofollow" – new ways to identify the nature of links — Search Central blog (10 September 2019)
15. Pagination, incremental page loading, and their impact on Google Search — Google Search Central
16. What is canonicalization — Google Search Central
17. Crawling December: CDNs and crawling — Search Central blog (24 December 2024)
18. Search Central Deep Dive Europe 2026: Apparently we're going to Barcelona — Search Central blog (6 July 2026)
19. Introduction to robots.txt — Google Search Central
20. No page information in search results — Google Search Console Help
21. Block Search indexing with noindex — Google Search Central
22. Build and submit a sitemap — Google Search Central
23. Help Google understand your ecommerce website structure — Google Search Central
24. How is a Google Sitemap different from an HTML sitemap? — Search Central blog (26 September 2005)
25. Temporarily pause or disable a website — Google Search Central
26. Redirects and Google Search — Google Search Central
27. Robots meta tag, data-nosnippet, and X-Robots-Tag specifications — Google Search Central
28. meta tags and attributes that Google supports — Google Search Central
29. Link best practices for Google — Google Search Central
30. Understand the JavaScript SEO basics — Google Search Central
31. Tell Google about localized versions of your page — Google Search Central
32. Control your snippets in search results — Google Search Central
33. AI features and your website — Google Search Central
34. Discover and your website — Google Search Central
35. Translated results in Google Search — Google Search Central
36. Search Engine Optimization (SEO) Starter Guide — Google Search Central
37. Robots Refresher: Future-proof Robots Exclusion Protocol — Search Central blog (28 March 2025)

38. Formalizing the Robots Exclusion Protocol Specification — Search Central blog (1 July 2019)
39. Helping you break the language barrier — Search Central blog (14 October 2008)
40. Translate — The Chromium Projects
41. Fix Search-related JavaScript problems — Google Search Central
42. Dynamic rendering as a workaround — Google Search Central
43. Fix lazy-loaded content — Google Search Central
44. URL Inspection tool — Google Search Console Help
45. List of Google user-triggered fetchers — Google
46. Crawling December: The how and why of Googlebot crawling — Search Central blog (3 December 2024)
47. Inside Googlebot: demystifying crawling, fetching, and the bytes we process — Search Central blog (31 March 2026)
48. Mobile site and mobile-first indexing best practices — Google Search Central
49. March 2023 Google SEO office hours — Google Search Central (SEO office hours transcript, March 2023)
50. December 2023 Google SEO office hours — Google Search Central (SEO office hours transcript, December 2023)
51. New robots tag: indexifembedded — Search Central blog (21 January 2022)
52. Verify your site ownership — Google Search Console Help
53. How to specify a canonical URL with rel="canonical" and other methods — Google Search Central
54. Fix canonicalization issues — Google Search Central
55. How to move a site — Google Search Central
56. Managing multi-regional and multilingual sites — Google Search Central
57. 5 common mistakes with rel=canonical — Search Central blog (8 April 2013)
58. General Guidelines — Google Search Quality Rater Guidelines (PDF, 11 September 2025)
59. Long context — Google AI for Developers (Gemini API docs)
60. Understand and count tokens — Google AI for Developers (Gemini API docs)
61. Specify your canonical — Search Central blog (12 February 2009)
62. Introduction to structured data markup in Google Search — Google Search Central
63. General structured data guidelines — Google Search Central
64. Structured data markup that Google Search supports — Google Search Central
65. Review snippet (Review, AggregateRating) structured data — Google Search Central
66. Event (Event) structured data — Google Search Central
67. Organization (Organization) structured data — Google Search Central
68. Loyalty program (MemberProgram) structured data — Google Search Central
69. Merchant shipping policy (ShippingService) structured data — Google Search Central
70. Merchant listing (Product, Offer) structured data — Google Search Central
71. Latest documentation updates — Google Search Central
72. Simplifying the search results page — Search Central blog (12 June 2025)
73. Rakuten increased time on site 1.5x with the recipe search experience on Google — Google Search Central case study (8 May 2018)
74. Introducing schema.org: Search engines come together for a richer web — Search Central blog (2 June 2011)
75. Rich Results Test — Google Search Console Help
76. Google Search's guidance on using generative AI content on your website — Google Search Central
77. Google image SEO best practices — Google Search Central
78. Image sitemaps — Google Search Central
79. Supporting AVIF in Google Search — Search Central blog (30 August 2024)
80. Remove images hosted on your site from search results — Google Search Central
81. Video SEO best practices — Google Search Central

82. Video sitemaps and alternatives — Google Search Central
83. How Google crawls locale-adaptive pages — Google Search Central
84. Spam policies for Google web search — Google Search Central
85. How search results may differ based on accented characters and interface languages — Search Central blog (1 September 2006)
86. How Google knows what language to show in search results — Google Search Help
87. 75% of Multilingual Websites Have Hreflang Implementation Mistakes — Search Engine Journal (15 February 2017)
88. Page indexing report — Google Search Console Help
89. Guidelines for sites with explicit content — Google Search Central
90. How we fought spam on Google Search in 2022 — Search Central blog (11 April 2023)
91. How we fought Search spam on Google in 2021 — Search Central blog (21 April 2022)
92. Google Search spam updates and your site — Google Search Central
93. Content policies for Google Search — Google Search Help
94. How Google Search organizes information — Google Search (How Search Works)
95. How AI powers great search results — Google blog (3 February 2022)
96. Expanding AI Overviews and introducing AI Mode — Google blog (5 March 2025)
97. Get started with Google Trends — Google Search Central
98. Google Trends — Google Trends
99. FAQ about Google Trends data — Google Trends Help
100. Compare search terms and topics — Google Trends Help
101. Compare Trends search terms — Google Trends Help
102. Find related searches — Google Trends Help
103. Understanding How the Google Trends Explore Page uses Gemini models — Google Trends Help
104. Compare search interest over time with quick comparisons — Google Trends Help
105. Explore the searches that are Trending now — Google Trends Help
106. An easier way to explore Search trends with Gemini — Google blog (14 January 2026)
107. 3 Trending Now updates to help you keep up with the latest trends — Google blog (14 August 2024)
108. AI, personalization and the future of shopping — Google blog (3 March 2025)
109. Learn all about Google Trends with our new YouTube tutorials — Search Central blog (12 September 2024)
110. Basics of Google Trends — Google News Initiative
111. Advanced Google Trends — Google News Initiative
112. Image editing in Gemini just got a major upgrade — Google blog (26 August 2025)
113. Google Trends Gains Explore Maps & Regional Breakdown — Search Engine Roundtable (19 August 2026)
114. Google Trends Explore Page Adds Category Filtering — Search Engine Roundtable (2 September 2026)
115. Use valid HTML to specify page metadata — Google Search Central
116. Analyze your social and video platform content performance in Search Console — Google Search Central
117. How to write and submit a robots.txt file — Google Search Central
118. Video indexing report — Google Search Console Help
119. Image metadata in Google Images — Google Search Central
120. October 2023 Spam Update — Search Central blog (4 October 2023)
121. Site names on Google Search now available for all languages — Search Central blog (7 September 2023)
122. How Google Search handles multilingual searches — Search Central blog (8 September 2023)

Search Central Live is a Google event. This is an independent attendee guide by Ibrahim Anjro, not an official Google publication. Statistics are as presented on stage, with their scope as stated.